

# Proxmox Storage DRS -- Internals

How the tool works, for whoever changes it

---

Bernd Zeimetz [bernd@bzed.de](mailto:bernd@bzed.de)

2026-09-07

Rendered from docs/internals/\*.md, SHA-256 ccbc9410ca15f8ccf4918e6ae26bfd0aadb25e554aa3cc292656ce90d978c868  
(sha256sum --check docs/internals.pdf.sha256 in the repository must succeed).

Copyright © 2026 Bernd Zeimetz. Licensed under the GNU Affero General Public License, version 3 or later.

## Contents

|                                                                                          |           |
|------------------------------------------------------------------------------------------|-----------|
| The pipeline, as specified and as built . . . . .                                        | 5         |
| Module layout, as it stands . . . . .                                                    | 6         |
| Why config.py depends on forecast.py . . . . .                                           | 7         |
| Where to read next . . . . .                                                             | 8         |
| <b>How a YAML file becomes a validated Config</b>                                        | <b>8</b>  |
| Resolution order . . . . .                                                               | 8         |
| The two-stage validation . . . . .                                                       | 9         |
| Secrets from the environment . . . . .                                                   | 9         |
| What is deliberately <i>not</i> checked here . . . . .                                   | 9         |
| ResolvedConfig: what the caller actually gets . . . . .                                  | 9         |
| <b>Persistent state: state.json</b>                                                      | <b>9</b>  |
| One dataclass tree, matching section 11.2's JSON exactly . . . . .                       | 10        |
| Reading degrades; writing does not . . . . .                                             | 10        |
| Locking: flock() is the mechanism, the JSON lock field is a label . . . . .              | 10        |
| What today's wiring into cli.py actually changes . . . . .                               | 11        |
| Cooldowns: read by topology.py and heuristic.py, not by this module . . . . .            | 11        |
| Locking and writing: apply's own responsibility . . . . .                                | 11        |
| Deliberately not implemented in this pass . . . . .                                      | 12        |
| <b>Forecasting: the protocol and its three implementations</b>                           | <b>12</b> |
| The one rule, in one place . . . . .                                                     | 12        |
| Why the upper bound, never the point estimate . . . . .                                  | 12        |
| storage_upper_bound(): section 10.1's per-disk sum, not a per-storage forecast . . . . . | 12        |
| QuantileForecaster . . . . .                                                             | 13        |
| SeasonalNaiveForecaster . . . . .                                                        | 13        |
| HoltWintersForecaster . . . . .                                                          | 13        |
| The section 10.2 backtest validation gate (phase 9) . . . . .                            | 13        |
| build_forecaster() . . . . .                                                             | 14        |
| <b>The Prometheus client and verify-metrics</b>                                          | <b>14</b> |
| _SessionLike / _ResponseLike: why not requests.Session directly . . . . .                | 14        |
| _get(): one request path, three endpoint shapes . . . . .                                | 14        |
| PromQL construction . . . . .                                                            | 15        |
| verify_metrics(): six checks, six functions . . . . .                                    | 15        |
| What cli.py does with the report . . . . .                                               | 15        |
| <b>Command dispatch, the mode-override rule, and why logs go to stderr</b>               | <b>15</b> |
| One parser, global options first . . . . .                                               | 15        |
| Command handlers: a dict, not a chain of ifs . . . . .                                   | 16        |
| --group: filtered once, right after build_topology() . . . . .                           | 16        |
| The --mode escalation rule . . . . .                                                     | 16        |
| Why logs go to stderr, not stdout . . . . .                                              | 16        |
| JsonFormatter: what ends up in one log line . . . . .                                    | 17        |
| --manual: preferring man(1), falling back only when necessary . . . . .                  | 17        |
| <b>The Proxmox VE API client</b>                                                         | <b>17</b> |
| Why proxmoxer . . . . .                                                                  | 17        |


|                                                                                              |           |
|----------------------------------------------------------------------------------------------|-----------|
| build_client(): the one place auth is assembled . . . . .                                    | 17        |
| PveClient: one method per section 3.5 endpoint . . . . .                                     | 17        |
| move_disk(): the one and only bytes/s -> KiB/s conversion . . . . .                          | 19        |
| <b>Building the disk/storage/group model . . . . .</b>                                       | <b>19</b> |
| D means every disk this module places, pinned or not . . . . .                               | 19        |
| One pass, in the order section 3.5 lists . . . . .                                           | 19        |
| The lock field comes from vm_config(), not a separate call . . . . .                         | 20        |
| Pin priority: _pin_reason() . . . . .                                                        | 20        |
| The per-disk cooldown pin . . . . .                                                          | 20        |
| Sizes: content is authoritative, config is the fallback . . . . .                            | 20        |
| U <sup>s e x t</sup> : everything not referenced . . . . .                                   | 21        |
| Section 11.4: /.../ storage patterns are expanded here, once . . . . .                       | 21        |
| reserve.py: one (C4)/(C5) evaluator, shared . . . . .                                        | 22        |
| <b>The load model . . . . .</b>                                                              | <b>22</b> |
| One function, one group, already-built input . . . . .                                       | 22        |
| Six queries, not six-times-groups queries . . . . .                                          | 22        |
| Coverage rejection excludes a disk from the group total, not just from its own l_d . . . . . | 22        |
| tpmstate0/unusedN are never rejected; efidisk0 is . . . . .                                  | 23        |
| T_g = 0 is “idle”, computed only from accepted disks . . . . .                               | 23        |
| L_s/u_s are the <i>current</i> assignment, not a candidate one . . . . .                     | 23        |
| compute_disk_load_series(): the same blend, as a time series . . . . .                       | 23        |
| What is still missing from phase 3 . . . . .                                                 | 24        |
| <b>Gating: deciding whether to act at all . . . . .</b>                                      | <b>24</b> |
| One function, three gates, in the order section 6 lists them . . . . .                       | 24        |
| last_load is real now, for show-load and plan . . . . .                                      | 24        |
| show-load’s gate line is diagnostic, not a real decision . . . . .                           | 25        |
| Cooldowns are not this module’s job . . . . .                                                | 25        |
| <b>The heuristic solver . . . . .</b>                                                        | <b>25</b> |
| evaluate_assignment() is the objective, kept separate from the search . . . . .              | 25        |
| Repair is unconditional, not weight-driven . . . . .                                         | 25        |
| Descend explores swaps, not only single moves . . . . .                                      | 26        |
| Descend also tries relocating a whole VM at once — a real gap, found by dogfooding . . . . . | 26        |
| The storage cooldown excludes a destination, never a source . . . . .                        | 27        |
| objective.spread_metric: two different quantities, not one rescaled . . . . .                | 27        |
| Proof this reproduces the plan, not just itself . . . . .                                    | 27        |
| What this pass deliberately does not do . . . . .                                            | 28        |
| <b>The MILP backends: CP-SAT and CBC . . . . .</b>                                           | <b>28</b> |
| Both backends solve, then hand off to heuristic.evaluate_assignment() . . . . .              | 28        |
| Lexicographic, not big-M — and why that is the only mode implemented . . . . .               | 28        |
| CP-SAT: section 5.5’s exact integer scaling . . . . .                                        | 29        |
| CBC: “direct transcription”, deliberately not scaled . . . . .                               | 29        |
| (C1)/(C3)/(C4)/(C5) are shared code, factored once per backend . . . . .                     | 30        |
| cli.py’s dispatch: auto cascades, an explicit backend still falls back . . . . .             | 30        |
| The storage cooldown: enforced in both stages, both backends (REVIEW.md S-04) . . . . .      | 30        |
| Deliberately not implemented in this pass . . . . .                                          | 31        |


|                                                                                           |           |
|-------------------------------------------------------------------------------------------|-----------|
| <b>Executing a plan: execute.py and its cli.py wiring</b>                                 | <b>32</b> |
| execute_plan() re-validates every move against the live cluster, not the plan . . . . .   | 32        |
| Why “done” needs three conditions, not one . . . . .                                      | 32        |
| VM locks are an open set, waited out, never whitelisted . . . . .                         | 33        |
| The injectable Clock . . . . .                                                            | 33        |
| Orphaned volumes: reported, never deleted . . . . .                                       | 33        |
| confirm’s callback lives in cli.py, not here . . . . .                                    | 34        |
| What execute_plan() deliberately does not do . . . . .                                    | 34        |
| cli.py: one planning pipeline, shared by plan and apply . . . . .                         | 34        |
| The payback gate: apply refuses on its own verdict (REVIEW.md S-02) . . . . .             | 35        |
| auto mode: timewindow.py and _run_auto_group() (phase 8) . . . . .                        | 35        |
| Concurrent execution: _execute_concurrent() and its non-blocking helpers . . . . .        | 37        |
| <b>Crash and two-instance recovery: crashrecovery.py</b>                                  | <b>38</b> |
| The two failure modes section 13 names, one mechanism . . . . .                           | 38        |
| parse_upid(): PVE’s UPID grammar, confirmed live . . . . .                                | 38        |
| /cluster/tasks’ own convention . . . . .                                                  | 38        |
| reconcile_inflight(): both halves . . . . .                                               | 39        |
| Writing before the crash: execute.py’s callbacks . . . . .                                | 39        |
| cli.py: the startup scan, before planning anything . . . . .                              | 40        |
| <b>Ordering the moves</b>                                                                 | <b>40</b> |
| One loop, reusing heuristic.py’s objective for the ratio . . . . .                        | 40        |
| cost_m is z_d — and that is not an approximation today . . . . .                          | 40        |
| The transient invariant, called with a single-move set always . . . . .                   | 40        |
| The heuristic can accept a residual violation; the scheduler cannot execute one . . . . . | 41        |
| final_assignment is the schedule’s own truth, not the heuristic’s target . . . . .        | 41        |
| What this pass deliberately does not do . . . . .                                         | 41        |
| Wired into plan, not yet into apply . . . . .                                             | 42        |
| <b>Migration cost and the payback test</b>                                                | <b>42</b> |
| Cost and benefit are computed from data other modules already have . . . . .              | 42        |
| headroom_src/headroom_dst: a plan formula this project cannot fill in . . . . .           | 42        |
| The reserve-override exemption . . . . .                                                  | 42        |
| The section 7.3 saturation guard: compute_move_cost()’s optional target . . . . .         | 43        |
| What this pass deliberately does not do . . . . .                                         | 43        |
| Wired into plan and apply alike, via _plan_group() . . . . .                              | 44        |

**What does this page answer?** Where does an invocation of pve-storage-drs go, which module owns which piece, and how much of IMPLEMENTATION\_PLAN.md's section 2 architecture actually exists right now?

## The pipeline, as specified and as built

IMPLEMENTATION\_PLAN.md section 2 describes seven stages: collect, join, gate, solve, cost, order, execute. As of this page, all seven exist, with **two** interchangeable stage-4 (solve) backends — the dependency-free heuristic and the MILP (CP-SAT/CBC) path — selected by solver.backend. Stage 7 (execute) is real for all three execution.mode values, including auto's own section 9.1 time-window/migration-count budgets, section 9.2 re-plan loop, and section 8.1/9.2's concurrent execution (execution.max\_concurrent\_migrations/max\_concurrent\_per\_storage above their default of 1, auto-only, strictly FIFO — see 92-execute.md's own "Concurrent execution" section). Do not take this page as a claim that every phase-8 knob does something — ../manual/30-safety-and-status.md is the authoritative per-command status: plan prints an ordered, transient-feasible, payback-checked plan (see ../manual/27-plan.md), and apply (see ../manual/28-apply.md) executes it in any mode, re-validating every move against the live cluster immediately before issuing it (section 9.2), waiting out a VM lock rather than failing (section 9.3.1), and not considering a move done until the task succeeded *and* the source volume is gone *and* the VM's lock is clear (section 9.3.2) — see 92-execute.md. state.json itself (state.py) is read by show-load and plan for drift history and cooldowns, and now **written** by apply after any run that actually executed a migration — see 15-state.md, 80-gates.md and 92-execute.md.





## Module layout, as it stands

| Module             | Responsibility                                                                                                                                                                                                                                                          | Plan section                                                      |
|--------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------|
| exceptions.py      | The DrsError hierarchy every deliberate failure raises                                                                                                                                                                                                                  | AGENTS.md section 5                                               |
| units.py           | Duration/size parsing ("24h", "200MiB") and human-readable formatting                                                                                                                                                                                                   | section 11                                                        |
| config.py          | Load /etc/pve/drs.yaml, jsonschema + semantic validation, the frozen dataclass config model                                                                                                                                                                             | section 11                                                        |
| config_schema.json | The jsonschema structural half of validation                                                                                                                                                                                                                            | section 11.1                                                      |
| state.py           | state.json: the load vector as of the last executed balance, cooldowns, the node-local advisory flock() — reading degrades, writing raises                                                                                                                              | section 11.2                                                      |
| forecast.py        | The Forecaster protocol, required_range_seconds, and quantile/seasonal_naive/holt_winters                                                                                                                                                                               | section 10                                                        |
| logging_setup.py   | Structured JSON logging to <b>stderr</b>                                                                                                                                                                                                                                | section 2.1 (amended, see <a href="#">40-cli-and-logging.md</a> ) |
| metrics.py         | PrometheusClient, PromQL construction, verify_metrics(), compute_disk_coverage()                                                                                                                                                                                        | sections 3.1-3.4                                                  |
| pve.py             | PveClient (built on proxmoxer), build_client()                                                                                                                                                                                                                          | section 3.5                                                       |
| topology.py        | build_topology(): the disk/storage/group join, D, S, U <sup>s e x t</sup> , (C2) pins                                                                                                                                                                                   | sections 3.5-3.7, 5.1, 5.3 (C2)                                   |
| reserve.py         | compute_reserve_status(): (C4)/(C5), shared by show-load today and the solver later; transient_charge_ok(): section 8.1's transient invariant, generalized to any number of concurrent charges on one target, shared by schedule.py (model-based) and execute.py (live) | section 5.3 (C4)/(C5), section 8.1                                |
| loadmodel.py       | compute_group_load(): the raw-series-to-l_d blend, min_coverage rejection, current L <sub>s</sub> /u <sub>s</sub>                                                                                                                                                       | section 4                                                         |
| gates.py           | evaluate_group_gates(): reserve override, drift, imbalance — the act/no-act verdict, with reasoning                                                                                                                                                                     | section 6                                                         |

| Module        | Responsibility                                                                                                                                                                               | Plan section     |
|---------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
| heuristic.py  | run_heuristic(): seed/repair/descend, and evaluate_assignment(), the section 5.4 objective shared with the MILP path too                                                                     | sections 5.4/5.5 |
| optimize.py   | solve(): CP-SAT/CBC, section 5.3's constraints, the lexicographic two-stage reserve solve                                                                                                    | section 5.5      |
| schedule.py   | order_moves(): transient-feasible ordering of a target assignment's moves, deadlock reporting                                                                                                | section 8        |
| payback.py    | evaluate_plan_payback(): the cost/benefit acceptance test, with a reserve-override exemption mirroring gates.py's                                                                            | section 7        |
| execute.py    | execute_plan(): pre-flight re-check per move, VM-lock wait, move_disk, the three-condition completion criterion, orphan detection on failure, auto's own time-window/migration-count budgets | section 9        |
| timewindow.py | current_deadline(): is now (local time) inside a configured execution.time_windows entry, and when does it close                                                                             | section 9.1      |
| cli.py        | Argument parsing, command dispatch, --manual, the mode-override rule, show-load, verify-storages, plan, apply (including auto's own re-plan loop)                                            | section 11.3     |

Every phase 8 knob is now real: time windows, max\_migrations\_per\_run, the re-plan loop, and concurrent execution (execution.max\_concurrent\_migrations/max\_concurrent\_per\_storage above their default of 1, auto-only, strictly FIFO) — see [92-execute.md](#) and [../manual/28-apply.md](#). schedule.py itself still only ever produces a strictly-sequential *order* (see below); execute.py's concurrent executor uses that same order as a launch queue rather than needing schedule.py to reason about overlapping in-flight windows itself. optimize.py (phase 6) is done: plan/apply pick CP-SAT, CBC or the heuristic per solver.backend, and both MILP backends now enforce the storage cooldown identically to the heuristic (see [91-optimize.md](#)). Cooldowns and drift history are read *and written*: topology.py's (C2) per-disk pin and heuristic.py's per-storage target exclusion consume state.py's cooldown data, and apply writes a fresh one for every disk and **both** of a move's storages a run actually migrated (see [15-state.md](#), [60-topology.md](#), [90-heuristic.md](#) and [92-execute.md](#)). Within modules that do exist: heuristic step 4 “polish” and (C2) format-compatibility eligibility in heuristic.py (see [90-heuristic.md](#)); schedule.py itself still only ever produces a strictly-sequential order (not execute.py's own concurrent *execution* of it, which is real – see above – but reasoning about which pairs of moves could safely be *scheduled* to overlap, priority-2 ordering, and staging, per [95-schedule.md](#)); and the payback re-solve-and-retry loop in payback.py (see [96-payback.md](#) – the section 7.3 saturation check itself is implemented there now, mirroring-phase only).

### Why config.py depends on forecast.py

config.py's semantic validation (\_check\_forecast\_window) must reject a window.lookback too short for the configured forecast.model at startup (section 11.1) — that check needs forecast.required\_range\_seconds(). Since forecast.py in turn needs config.ForecastConfig for its type signatures, the import is deferred to inside the validation function rather than at module level, breaking what would otherwise be an import cycle. See the comment at the top of config.py's \_check\_forecast\_window if you touch this.

## Where to read next

- [10-configuration.md](#) — how a YAML file becomes a validated Config, and where every default actually lives.
- [15-state.md](#) — state.json’s dataclasses, why reading it degrades but writing raises, the real flock() lock and the rename-vs -in-place bug it takes to get that wrong, what show-load/plan actually get from it today, and how its cooldown data reaches topology.py/heuristic.py.
- [20-forecasting.md](#) — the forecaster protocol and its three implementations.
- [30-metrics.md](#) — the Prometheus client and the six verify-metrics checks.
- [40-cli-and-logging.md](#) — command dispatch, the --mode escalation rule, and why logs go to stderr.
- [50-pve-api.md](#) — the PVE API client, why it is built on proxmoxer, and two things verified against a real cluster.
- [60-topology.md](#) — the disk/storage/group join, the section 5.1.1 foreign-volume accounting, the shared (C4)/(C5) evaluator, and the per-disk cooldown pin.
- [70-loadmodel.md](#) — the section 4 raw-series-to-l\_d blend, min\_coverage rejection, and why tpmstate0/unusedN are exempt from it but efidisk0 is not.
- [80-gates.md](#) — the section 6 act/no-act verdict, how state.json’s drift history now reaches it, and why the per-disk/ per-storage cooldowns live in topology.py/heuristic.py instead.
- [90-heuristic.md](#) — the section 5.4/5.5 objective and the seed/repair/descend search, cross-checked against section 14’s exact objective totals, the storage-cooldown destination filter and its repair-side exemption, and what “polish” and format eligibility still owe.
- [91-optimize.md](#) — the CP-SAT and CBC MILP backends, section 5.5’s exact integer coefficient scaling (and the  $\gamma$ -term trap it exists to avoid), why the reserve is solved in two lexicographic stages rather than one big-M objective, and cli.py’s solver.backend dispatch (auto’s cascade, and why an explicitly forced backend still falls back to the heuristic rather than failing).
- [95-schedule.md](#) — ordering a target assignment’s moves under the section 8 transient invariant, why cost\_m = z\_d is exact today (not an approximation), and why the heuristic accepting a residual violation can still mean the scheduler reports a deadlock.
- [96-payback.md](#) — the section 7 cost/benefit test, the headroom\_src/headroom\_dst gap in the plan’s own cost formula, and why a reserve-fixing plan always passes the economic test.
- [92-execute.md](#) — why a completed move needs three conditions, not one, why a VM lock is waited out rather than whitelisted, the injectable Clock, how cli.py’s \_plan\_group() gives plan and apply one shared pipeline instead of two, and auto mode’s time-window budget and re-plan loop (timewindow.py, cli.\_run\_auto\_group()).
- [93-crashrecovery.md](#) — section 13’s startup scan for a move\_disk left running by a crash or a second instance, the cluster-wide /cluster/tasks read that state.json’s own flock() cannot do, and how a UPID reaches disk *before* the move that owns it can crash the engine.

## How a YAML file becomes a validated Config

**What does this page answer?** What exactly happens between pve-storage-drs -c foo.yaml plan and a type-checked Config object, and where does each of the section 11.1 validation rules actually live? Describes proxmox\_storage\_drs/config.py and config\_schema.json.

## Resolution order

resolve\_config\_path() implements IMPLEMENTATION\_PLAN.md section 11’s three-source order: --config (cli\_path), then \$PVE\_STORAGE\_DRS\_CONFIG, then config.DEFAULT\_CONFIG\_PATH (/etc/pve/drs.yaml). It also returns was\_explicit, which load\_config() uses to decide the error message on a read failure: an explicitly named path that cannot be read is always fatal with no fallback, per section 11, but the message additionally points at config/drs.example.yaml only when the failing path was the unnamed default — there is nothing useful to suggest when the operator named the path themselves.

## The two-stage validation

1. `_validate_schema()` loads `config_schema.json` via `importlib.resources` (it ships as package data — see `pyproject.toml`'s `[tool.setuptools.package-data]`) and runs it through `jsonschema`. Every object in the schema sets `additionalProperties: false`; a typo'd key is caught here as a validation error, not silently ignored, which is what keeps section 15.1's knob-to-formula table trustworthy (a key with no formula would otherwise go unnoticed).
2. `_build_config()` walks the now-schema-valid raw dict into the frozen dataclasses defined earlier in the module, applying every default inline via `.get(key, default)`. These defaults are also what `config/drs.example.yaml` documents and what the manual's [../manual/10-configuration.md](#) describes — there is deliberately no second copy of a default anywhere else in the codebase.
3. `_validate_semantics()` runs the section 11.1 rules `jsonschema` cannot express — cross-field comparisons, or rules needing a value derived from two independently-optional settings. Each rule is a small `_check_*` function appending to a shared errors/warnings list (kept separate functions rather than one large one, partly for flake8's max-complexity and partly because each is independently testable). errors become one `ConfigError` naming every problem found, not just the first; warnings are returned to the caller (`ResolvedConfig.warnings`) for `cli.py` to log — section 11.1's `saturation_load` check is the one warning-only rule today.

Every duration or size field is parsed **once**, at this point, via `units.parse_duration_seconds/parse_size_bytes`, and stored on the dataclass already in its canonical unit (seconds, bytes) — nothing downstream re-parses a string.

## Secrets from the environment

`_build_config()` fills `proxmox.auth.password/token_secret` from `PVE_PASSWORD/PVE_TOKEN_SECRET` **only when the file's own value is null** — a secret actually written in the file is used as-is, never silently overridden by an environment variable a deploy script forgot to unset. This one-directional fallback is what section 11's "keep secrets out of `/etc/pve`" guidance is for: the file can omit the field entirely and rely on the environment, or set it explicitly and take responsibility for the file's own permissions.

## What is deliberately *not* checked here

Two section 11.1 rules need live cluster data this module does not have and are therefore **not** part of `config.py`:

- storage ids actually existing in the cluster;
- `execution.source_release.timeout` / `gates.cooldown_per_storage` against a storage's real `saferemove_throughput`.

Both belong to `pve.py/topology.py` (which do the live cluster fetch) and `pve-storage-drs verify-storages` — see [../manual/30-safety-and-status.md](#) for current status.

## ResolvedConfig: what the caller actually gets

`load_config()` returns a `ResolvedConfig`, not a bare `Config`: the resolved path and the file's sha256 travel with it, because `IMPLEMENTATION_PLAN.md` section 11 requires every run to log both — a plan must always be traceable to the exact file that produced it, even though the file is never re-read after startup.

## Persistent state: `state.json`

**What does this page answer?** What does `state.py` actually store, how does reading it degrade instead of failing, why is the advisory lock a real `flock()` rather than a JSON field, and what does today's wiring into `show-load/plan` actually change? Describes `proxmox_storage_drs/state.py`.

### One dataclass tree, matching section 11.2's JSON exactly

State mirrors the plan's own example document field for field: lock (LockInfo | None), last\_balance (at plus a "<group>:<vmid>:<device>" -> l\_d load\_vector), cooldowns (disk/storage, each a key -> timestamp mapping), inflight\_upids, staged\_disks. disk\_state\_key()/ storage\_state\_key() build the two key shapes section 11.2 specifies — group-prefixed specifically so a vmid reused after a VM is destroyed and recreated in a *different* group cannot collide with its old entry. empty\_state() is what a fresh install, or a lost file, degrades to: no lock, no balance, no cooldowns, nothing in flight.

### Reading degrades; writing does not

load\_state() never raises. A missing file is the ordinary first-run case (silent); a present-but-corrupt, empty, or wrong-schema\_version one is still just empty\_state(), but logged at WARNING — section 11.2's own words are “losing this file is safe but not free: cooldowns and drift history reset”, not “state.json existing and being readable is required to run”. This mirrors show-load's choice to degrade a group's load to “unavailable” on a Prometheus outage rather than fail the whole command — neither Prometheus nor state.json is an *explicitly requested* input the way -c/--config PATH is, and only the latter is a hard failure on error (.agents/domain-invariants.md section 9).

save\_state\_atomic(), by contrast, raises StateError on any failure — writing (or locking) is something a caller actively asked this module to do, not a read with a documented, safe fallback. It writes to a temp file in the same directory, fsyncs, then os.replace()s over the real path, so a concurrent load\_state() can never observe a half-written file — which is exactly what lets every read in this codebase skip locking entirely (next section).

### Locking: flock() is the mechanism, the JSON lock field is a label

acquire\_lock() takes a real, non-blockingfcntl.flock(LOCK\_EX | LOCK\_NB) on the state file itself. That is the actual, kernel-enforced mutual exclusion section 11.2 asks for, and it is correct by construction for two instances on the *same host* — .agents/domain-invariants.md section 9 is explicit that “the fcntl lock cannot see another node”; cross-node coordination is the separate in-flight-UPID scan (section 13), not this module's job.

The lock: {pid, host, acquired\_at} field recorded *inside* the JSON is **descriptive metadata for a human reading the file, not the exclusion mechanism** — by the time this module's code overwrites it, the OS has already granted exclusive ownership, so whatever a previous run last wrote there is necessarily stale and is unconditionally replaced, never inspected to decide whether to proceed. This is also how section 11.2's “if lock.pid is not alive on lock.host, reclaim it” is satisfied, without a separate liveness check to get subtly wrong: a process that died released its flock() the moment its file descriptor closed (any exit, including a crash), so the *next* acquire\_lock() call simply succeeds. acquire\_lock() returns None — not an exception — when another live instance already holds it, matching section 11.2's “a live PID means another instance is running: exit 0 quietly”; deciding what “quietly” means, and which exit code, is left to the caller.

**A real bug this design surfaced during testing, worth remembering if you touch this code:** the first implementation had acquire\_lock() write its lock metadata via save\_state\_atomic() — the same temp-file-plus-rename used everywhere else. That silently breaks the lock: a rename replaces the directory entry with a *new* inode that was never flock()'d, so a second acquire\_lock() opens that new inode fresh and locks it with no conflict at all — test\_a\_second\_acquire\_while\_the\_first\_is\_held\_returns\_none caught this immediately. The fix, \_write\_state\_to\_locked\_fd(), writes **in place** into the already-open, already-locked file descriptor (lseek to 0, write, ftruncate to the new length, fsync) — never a rename — for as long as the lock is held.

## What today's wiring into cli.py actually changes

show-load and plan both call `_last_loads_by_group()` once, right after reading configuration, and use the result for both `loadmodel.compute_group_load()`'s `last_known_loads` and `gates.evaluate_group_gates()`'s `last_load` — one state read per run, one function computing the per-group slice, not two independent implementations (AGENTS.md section 5). Neither command takes the advisory lock: both are read-only reports that never execute a migration, and taking the exclusive lock for one would make an in-progress apply block plan/show-load for no safety reason the plan text supports. `apply` (docs/internals/92-execute.md) is the one command that does take it — before touching PVE or Prometheus at all, so a second concurrent instance exits quietly without paying for either round-trip first.

The practical effect: a `state.json` with a real `last_balance` now makes the drift gate genuinely suppress an ACT verdict the imbalance alone would otherwise trigger — see `test_show_load_gate_reflects_real_drift_history_from_state_json` and `test_plan_gate_also_reflects_real_drift_history_from_state_json` in `tests/unit/test_cli.py` for the exact before/after. Without a `state.json` on disk (the common case until `apply` has actually executed a migration), behaviour is unchanged from before this module existed: `last_load=None`, drift gate skipped, “reserve override, else imbalance”.

## Cooldowns: read by topology.py and heuristic.py, not by this module

`state.py` only stores and queries cooldown timestamps (`active_disk_cooldowns()`/`active_storage_cooldowns()`); it does not decide what a cooldown *means* to the solver. `build_topology()` (via a new `state/now` parameter, both defaulting to “no cooldowns”/“real clock”) computes each group's active disk cooldowns once and feeds them into `topology._pin_reason()` — section 5.3 (C2)'s “d is within its per-disk cooldown -> also pin to current”, reported as `cooldown: moved recently, <time> left on gates.cooldown_per_disk`, in the plan's own priority order (after a snapshot pin, before a lock pin). `cli.py` separately computes each group's active *storage* cooldowns and passes them to `heuristic.run_heuristic()`, which excludes them as a move/swap **destination** in `_descend()` only — see docs/internals/60-topology.md and docs/internals/90-heuristic.md for the two halves in detail, including the deliberate asymmetry: a live (C4)/(C5) repair move ignores the storage cooldown entirely (section 13's reserve-override principle), while nothing yet exempts a disk-cooldown pin from blocking a repair the same way — a known, documented limitation, not an oversight.

`apply` is what actually calls `with_recorded_cooldown()` now, once per group, for every disk/storage a run's `execute_plan()` call actually migrated — a fresh disk cooldown at the disk's *new* location, and a fresh storage cooldown for the move's **destination** only, mirroring `_descend()`'s own asymmetry above (docs/internals/92-execute.md). A group `apply` never acts on (the gate said NO ACTION, or nothing executed before a `load_error/replan_needed` stopped it) leaves that group's cooldowns untouched, exactly like before this was wired in.

## Locking and writing: apply's own responsibility

`acquire_lock()/release_lock()/save_locked_state()` are exercised now: `apply` is the one command that can execute a migration, so it is the one this lock protects (see docs/internals/92-execute.md). It acquires the lock first, builds up the run's State in memory as each group finishes executing (`with_recorded_balance()/with_recorded_cooldown()`), and writes it once at the very end via `save_locked_state()` — **never** `save_state_atomic()`'s `rename`, which would silently detach the very lock just taken (see this page's own account of that bug, above) — before `release_lock()` in a `finally`, so a mid-run exception never leaves the lock held. `show-load/plan` still never take it, for the same read-only-report reason as always.

## Deliberately not implemented in this pass

- **A disk-cooldown pin is not exempted for an active reserve violation the way the storage cooldown is.** `topology.py` decides a disk's pin before the group's reserve. `ReserveStatus` is even computable (it needs every disk in the group, which is still being built), so there is no cheap way today to ask “would un-pinning this disk help fix a live violation?” at the point the pin decision is made. The practical effect is bounded and safe, never silently wrong: `heuristic._repair()` already reports an unresolvable residual violation via `reserve_penalty_term` whenever no movable disk can fix it (the same path `test_repair_does_not_oscillate_when_no_target_can_fully_absorb_the_violation` exercises) — a disk-cooldown pin can only ever add to the set of cases that hit that path, never bypass it. Revisiting this needs restructuring `topology.py` into two passes (build every disk first, compute reserve status per storage, then finalize cooldown pins), which is a real, separately-scoped change.
- **inflight\_upids** is now written and read — see `docs/internals/93-crashrecovery.md` for the full mechanism (`with_inflight_upid()/without_inflight_upid()`, `crashrecovery.py`'s startup scan, and `cli.py`'s synchronous-write callback into `execute_plan()`).
- **staged\_disks** exists in the dataclass and round-trips correctly, but nothing writes or reads it yet — it belongs to `schedule.py`'s staging (section 8.3 option 1), not implemented.

## Forecasting: the protocol and its three implementations

**What does this page answer?** How does `forecast.py` decide how much history a model needs, and what does each of `quantile/seasonal_naive/holt_winters` actually compute? Describes `proxmox_storage_drs/forecast.py`.

### The one rule, in one place

`IMPLEMENTATION_PLAN.md` section 10.1's table says how much history each model needs, independent of `window.lookback` (the *decision* window). That rule has exactly one implementation: three small pure functions, `_quantile_required_range_seconds`, `_seasonal_naive_required_range_seconds` and `_holt_winters_required_range_seconds`, each called from **two** places — the free function `required_range_seconds()` that `config.py`'s startup validation uses (before any `Forecaster` is constructed), and the matching `Forecaster.required_range()` method on the concrete class. Duplicating the arithmetic between those two call sites, rather than sharing the pure function, is exactly the kind of “second copy of a rule” `AGENTS.md` section 5 forbids — a change to the Holt-Winters requirement made in only one of them would silently desynchronize config validation from what the forecaster itself believes it needs.

### Why the upper bound, never the point estimate

Every `Forecaster.predict()` returns a `Forecast(point_estimate, upper_bound)`. Nothing in this module enforces that callers use `upper_bound` — that discipline belongs to the caller: the section 7.3 saturation guard (`payback.py`'s `compute_move_cost()`, via `storage_upper_bound()` below) already only ever reads `.upper_bound`; the optimizer does not consume a forecast at all yet. The asymmetry is stated in the module docstring and repeated here because it is easy to get backwards under time pressure: overestimating costs a slightly worse balance, underestimating risks scheduling a mirror onto a storage that is about to saturate.

### storage\_upper\_bound(): section 10.1's per-disk sum, not a per-storage forecast

Section 10.1 is explicit that  $\hat{L}_s(\Delta)$  is  $\sum_{d : x_{d,s}=1} \hat{u}_d(\Delta)$  — every disk on  $s$  forecast **independently**, then summed — never the forecast of  $s$ 's own already-summed series. `storage_upper_bound()` is exactly that sum, given a `Forecaster`, `{disk_key: TimeSeries}` (typically `loadmodel.compute_disk_load_series()`'s own output), the disk keys currently on one storage, and a horizon. Summing upper bounds this way is deliberately conservative (it assumes every disk peaks together — the right direction for

a guard whose failure mode is starting a mirror onto an already-busy array), and is real, measurable extra conservatism whenever a group's disks do not actually peak in lockstep: forecasting one already-summed series directly would let one disk's trough offset another's peak, understating the storage's own worst case. A disk key with no fetched series at all forecasts as an empty one (0.0), never a `KeyError` — this run may simply have no history for a disk yet.

### QuantileForecaster

The default. `predict()` takes the `window.quantile/window.upper_quantile` percentiles of the raw series with no fitting at all. `_quantile()` implements linear-interpolation percentile matching `numpy.percentile`'s default method, by hand — deliberately not using `numpy`, since the tool must stay usable with only `requests + ruamel.yaml + jsonschema` installed (IMPLEMENTATION\_PLAN.md section 2.1). `horizon` is accepted (to satisfy the Forecaster protocol) and immediately discarded: the quantile model has no notion of a horizon-dependent forecast.

### SeasonalNaiveForecaster

Groups historical samples by hour-of-day (`int((ts // 3600) % 24)`) and takes the median as the point estimate, the configured upper quantile across the same bucket as the bound. `now_epoch_seconds` selects which hour-of-day bucket the *current* moment falls into; `predict()` ignores everything outside that bucket. A bucket with no samples returns `Forecast(0.0, 0.0)` rather than raising, matching IMPLEMENTATION\_PLAN.md section 4's rule for an idle group's raw quantities: no data is zero, not an error, at this layer (the caller decides whether zero is trustworthy).

### HoltWintersForecaster

`statsmodels` is an optional dependency (Suggests, not Depends — see `debian/control`), so it is imported **only inside `predict()`**, never at module level; `debian/tests/import-all` is what would catch a regression here. Two situations fall back to `QuantileForecaster`, both logged at warning level rather than failing silently:

1. fewer than  $2 * \text{seasonal\_periods}$  samples in the series (section 10.2's own rule — a fit on too little data is worse than no fit);
2. `statsmodels` is not importable at all.

The upper bound is `point_estimate + residual_z * stdev(in-sample residuals)` — not derived from the model's own confidence interval, which `statsmodels`'s `ExponentialSmoothing.fit()` does not expose directly for every configuration. `residual_z` (`forecast.holt_winters.residual_z`, default 2.0) is what an operator tunes if this bound turns out too tight or too loose in practice.

## The section 10.2 backtest validation gate (phase 9)

Fitting successfully (`HoltWintersForecaster`'s own two fallback conditions above) is not the same question as fitting *accurately* — section 10.2's backtest is what actually answers the second one, once per group, before `seasonal_naive/holt_winters` is trusted to drive the section 7.3 saturation guard at all. `quantile` is never backtested: it does no fitting, so there is nothing to validate and nothing more conservative to fall back to.

`backtest_error()` implements the plan's own recipe literally: fit the candidate forecaster on `[now-2W, now-W)`, `predict W` ahead (`W` is `window.lookback_seconds`, the same “decision window” every forecaster is already tied to for its point estimate), and compare that single point estimate against the *actual* mean observed over `[now-W, now]`. The result is a **relative** error — `|predicted - actual| / actual`, normalizing by `predicted` instead only when `actual` is exactly zero (a nonzero prediction against true silence then reads as a full, bounded miss, 1.0, rather than a division by zero; two zeros, correctly predicted silence, score a perfect 0.0) — so it is directly comparable against `gates.imbalance_threshold` (both

plain fractions in `[0, 1]`, section 10.2's own choice of yardstick). `backtest_error()` returns `None`, not `0.0` or an exception, when series does not cover a full 2w to split into two non-empty halves at all; `backtest_validated()` treats that exactly like a failed backtest — a fresh deployment with no track record yet does not get a free pass just because it has not been disproven, matching `HoltWintersForecaster`'s own existing “not enough samples yet → fall back” precedent rather than contradicting it.

**Validated once per group, against the group's own aggregate series, not once per disk.** `group_aggregate_series()` sums every disk's own series at the union of their timestamps (a disk missing a sample at some timestamp contributes 0, the identical “absence means no I/O” convention `loadmodel._blend_loads()` already uses). `forecast.model` is a single, deployment-wide configuration choice — it cannot differ disk by disk — so validating it once, cheaply, against one representative series is the right granularity; backtesting separately per disk per move would be both more expensive and would raise an unanswerable question (use the fancier model for the disks it happens to fit and quantile for the rest, *within the same run*?) that the plan itself never poses.

`cli._backtest_gated_forecaster()` is the one caller: built once per group inside `_saturation_forecast_inputs()`, right after `build_forecaster()` and `loadmodel.compute_disk_load_series()`, before either is used to derive any move's `l_hat_src/l_hat_dst`. A model that fails validation falls back to a fresh `QuantileForecaster` (routed through `build_forecaster()` again with `forecast.model` overridden to “quantile”, never hand-constructed — the one factory, see below), logged at warning with the model name that failed.

### **`build_forecaster()`**

The one factory function that turns a `ForecastConfig` plus the ambient window/metrics values into a live `Forecaster` instance. Nothing else in the codebase should construct a `QuantileForecaster/SeasonalNaiveForecaster/HoltWintersForecaster` directly once a caller exists that needs one from config — route it through here so the model-name-to-class mapping stays in one place.

## **The Prometheus client and verify-metrics**

**What does this page answer?** How does `metrics.py` talk to Prometheus without ever touching the network in a test, and what does each of the six `verify-metrics` checks actually query? Describes `proxmox_storage_drs/metrics.py`.

### **`_SessionLike / _ResponseLike`: why not `requests.Session` directly**

`PrometheusClient.__init__` accepts `session: _SessionLike | None`, a Protocol defined in this module with only the one `get(...)` method the client actually calls — not `requests.Session` itself. `.agents/testing.md` forbids any test talking to a real Prometheus; a structural protocol lets a test hand in a small dataclass with a matching `get()` instead of mocking or subclassing the real (network-capable) session. `PrometheusClient()` with no `session` argument still constructs a real `requests.Session()`, which satisfies the protocol structurally — nothing special is needed to make the real thing fit.

### **`_get()`: one request path, three endpoint shapes**

Every one of `instant_query`, `range_query` and `label_values` goes through the private `_get()`, which does the HTTP call, checks the status code, parses JSON, checks Prometheus's own status: “success” field, and returns `payload["data"]`. That field's *shape* genuinely differs by endpoint — a dict with a `result` list for `query/query_range`, a bare list for `label/<name>/values` — which is why `_get()` is typed `Any` rather than picking one shape; each public method knows which shape it asked for and narrows accordingly. Every failure mode (a transport error, a non-200 status, unparseable JSON, an unsuccessful Prometheus

response) raises `MetricsError` from exactly this one place, so a caller catching `MetricsError` around any of the three public methods sees every failure mode uniformly.

### PromQL construction

`build_rate_promql()` and `build_quantile_over_time_promql()` are pure string-building functions, deliberately factored out of any query-issuing code so they are unit-testable without a fake session at all — see `IMPLEMENTATION_PLAN.md` section 3.4 for the exact expressions they implement. `raw_metric_name()` is the one place that maps a `RAW_METRIC_FIELDS` entry ("`read_ops`", ...) to the configured metric name on a `MetricsConfig`, via `getattr` — this is what lets `verify_metrics()` iterate “every configured raw metric” without hand-listing the six field names a second time anywhere in this module.

### `verify_metrics()`: six checks, six functions

Each `_check_*` function implements exactly one `IMPLEMENTATION_PLAN.md` section 3.3 numbered check and returns `Findings` (plus, where relevant, the data the report carries forward):

| Function                                   | Section 3.3 step       | Returns                                                             |
|--------------------------------------------|------------------------|---------------------------------------------------------------------|
| <code>_check_metric_names_exist</code>     | 1                      | findings only                                                       |
| <code>_check_sample_series</code>          | 2 + 3 (labels present) | findings, { <code>metric_name</code> : <code>sample_labels</code> } |
| <code>_check_device_label_collision</code> | 4                      | one <code>Finding</code> or <code>None</code> (pure, no I/O)        |
| <code>_check_coverage</code>               | 5                      | findings, { <code>DiskKey</code> : <code>coverage_fraction</code> } |
| <code>_check_observed_spacing</code>       | 6                      | findings, <code>observed_spacing_seconds</code>                     |

`_check_coverage` and `_check_observed_spacing` both use `metrics.read_ops` as the one representative metric rather than probing all six: a coverage or spacing gap is a property of the underlying Telegraf scrape, not of which of the six raw quantities is read, so probing all six would be six times the Prometheus load for no additional information.

`VerifyMetricsReport.ok` is `True` iff no `Finding` has `level == "error"` — `cli.py`’s `verify-metrics` handler uses exactly this property to decide the process exit code, so there is one definition of “the verification passed.”

### What `cli.py` does with the report

`_handle_verify_metrics` in `cli.py` constructs a `PrometheusClient` from `resolved.config.prometheus`, calls `verify_metrics()`, and renders either the human form (`[level]` message lines) or, under `--json`, a JSON-serializable dict — `DiskKey` keys are turned into "`vmid:device`" strings there, since JSON object keys must be strings and this module keeps no opinion about output formatting.

### Command dispatch, the mode-override rule, and why logs go to `stderr`

**What does this page answer?** How does `cli.py` turn `argv` into a dispatched command, what exactly does overriding `--mode log`, and why does structured logging go to `stderr` instead of the `stdout` the plan originally specified? Describes `proxmox_storage_drs/cli.py` and `proxmox_storage_drs/logging_setup.py`.

### One parser, global options first

`build_parser()` defines every global option (`-c/--config`, `--group`, `--mode`, `--json`, `-v/--quiet`, `--version`, `--manual`) on the **top-level** `ArgumentParser`, not on the subparsers, which is what makes `argparse` itself enforce `IMPLEMENTATION_PLAN.md` section 11.3’s “global options before the subcommand” for free — there is no hand-written ordering check anywhere. `--help`’s defaults come from the real

default= values passed to `add_argument`, combined with `argparse.ArgumentDefaultsHelpFormatter`; there is deliberately no second, hand-maintained list of options and defaults anywhere in this module, the manpage, or the manual (AGENTS.md section 8.5).

`main()` handles `--version` and `--manual/help` **before** requiring or loading any configuration — printing the version or the manual page must never depend on `/etc/pve/drs.yaml` existing. Only once a real subcommand is present does it call `config.load_config()`.

### Command handlers: a dict, not a chain of ifs

`_COMMAND_HANDLERS`: `dict[str, CommandHandler]` maps each subcommand name to a function of `(ResolvedConfig, argparse.Namespace, effective_mode) -> int`. Every subcommand not yet implemented (`apply`, `explain` — see `./manual/30-safety-and-status.md` for which ones that currently is) gets a handler from `_make_not_yet_implemented_handler(name)`, a closure factory rather than a lambda with a default-argument trick, because mypy's strict mode cannot infer a bare lambda's parameter types cleanly here — see the comment at that function if you are tempted to simplify it back to a one-liner. `_COMMAND_HANDLERS["verify-metrics"]` is overwritten with the real `_handle_verify_metrics` once `metrics.py` exists to back it. Adding a new implemented command is exactly this: write the handler, assign it into the dict, done — `main()`'s dispatch does not change.

### --group: filtered once, right after build\_topology()

`_filter_groups(topology, args.group)` is the one place `--group` is actually read. `show-load`, `verify-storages` and `plan` each call it immediately after `build_topology()`, replacing that `Topology`'s groups tuple with the (still topology-order) subset named — every render function downstream just iterates `topology.groups` as before and needs no `--group` awareness of its own. `verify-metrics` does not call it: that command validates configured metric/label names against Prometheus directly and never iterates groups at all, so there is nothing for `--group` to restrict there. A name that matches no configured group raises `DrsError` rather than silently producing an empty report — the same “an explicitly named thing that doesn't resolve is a hard failure” rule `config.load_config()` already applies to `-c/--config PATH` (REVIEW.md R-03: `--group` was previously parsed and documented, but no handler ever read `args.group` at all).

### The --mode escalation rule

`apply_mode_override(configured_mode, override)` is the entire implementation of section 11.3's rule: moving *down* the ordering `dry-run < confirm < auto` (toward more caution) logs at `INFO`; moving *up* it (toward less caution — removing a barrier the operator's own config set) logs at `WARNING`, naming both values. This is deliberately a small pure-ish function (its only side effect is the log call) rather than inlined into `main()`, so it is unit-testable against a `caplog` fixture without constructing a full CLI invocation.

### Why logs go to stderr, not stdout

IMPLEMENTATION\_PLAN.md section 2.1 originally specified `stdout` for structured JSON logs. This was amended in the same commit that added `logging_setup.py`: `--json` (section 9.5) emits the plan report on **stdout**, and interleaving log lines with that report on the same stream would make it unparseable by anything downstream. Logging instead goes to **stderr**; a systemd service unit captures both streams into the same journal regardless, so nothing about “captured by journald” is lost. This is the one place in the codebase where the implementation and IMPLEMENTATION\_PLAN.md disagree with the *original* plan text on purpose — AGENTS.md section 7 rule 2 is why the plan text itself was corrected in the same commit rather than left to silently diverge from the code.

`configure_logging()` is called exactly once, from `main()`, after the `--version/--manual` early exits (which must not touch logging configuration at all) and before any config-dependent code runs, so that a config-loading failure is itself logged consistently with everything after it.

**JsonFormatter: what ends up in one log line**

Every field on a `logging.LogRecord` that is *not* one of the standard attributes (computed once, at import time, into `_STANDARD_RECORD_ATTRS`) is folded into the JSON payload — this is what lets any module call `logger.info(..., extra={"event": "gate_decision", "threshold": 0.2})` and have event/threshold appear as top-level JSON keys with no formatter changes required. `sort_keys=True` keeps output byte-stable for anything that diffs log lines in a test or a log-aggregation pipeline.

**--manual: preferring man(1), falling back only when necessary**

`show_manual()` tries `man pve-storage-drs` first, via `subprocess.run` rather than `os.execvp`: `man(1)` itself detects a non-tty stdout and disables its pager automatically, which is what satisfies “never answer with a URL alone” (AGENTS.md section 8.5) whether the invocation is interactive or piped — there is no need for this module to duplicate that tty detection. `_fallback_manual_text()` is reached only when `man(1)` itself is missing or has no entry (an uninstalled source checkout, or a minimal system with no `man-db`): it walks up from `cli.py`’s own file location looking for `man/pve-storage-drs.1.md` in a source tree, and only if that also fails prints a short message naming the real installed-package path and the in-tree file — a local, actionable path, never a bare URL.

**The Proxmox VE API client**

**What does this page answer?** Why is `pve.py` built on `proxmoxer` instead of a hand-rolled ticket/CSRF client, and what does each of its methods actually call? Describes `proxmox_storage_drs/pve.py`.

**Why proxmoxer**

IMPLEMENTATION\_PLAN.md section 3.5 explains the decision in full; the short version is `proxmoxer`’s **backend abstraction**. The same `ProxmoxAPI` attribute-chaining interface (`api.nodes(node).qemu(vmid).config.get()`) is available over plain HTTPS (what `build_client()` uses today) or over SSH, either shelling out to the system’s own `ssh+pvesh` (`openssh`) or with an in-process client (`ssh_paramiko`). A deployment that cannot open the API port to the management host becomes a different keyword argument in `build_client()`, with no change anywhere else — not to `PveClient`’s methods, and not to any of their callers.

**build\_client(): the one place auth is assembled**

`config.py`’s `ProxmoxConfig` keeps `verify_ssl` (bool) and `ca_file` (path-or-None) as two separate knobs, because that is how an operator thinks about them; requests (and so `proxmoxer`’s https backend) wants one value, a bool or a CA-bundle path, passed as `verify`. `build_client()` is the one place that reconciles the two. It is also the one place `proxmox.auth.token_id`’s `user@realm!tokenname` format is split into `proxmoxer`’s separate `user/token_name` keyword arguments — never duplicate that parsing anywhere else.

Every failure `ProxmoxAPI(...)` itself can raise (`proxmoxer.AuthenticationError`, or a `requests.RequestException` if the host is simply unreachable) is caught here and re-raised as `PveApiError`, so every caller of `build_client()` and every `PveClient` method has exactly one exception type to catch.

**PveClient: one method per section 3.5 endpoint**

Every method is a single API call, wrapped through the private `_call()` helper, which is the one place `proxmoxer.ResourceException` (HTTP-level API errors) and `requests.RequestException` (transport failures `proxmoxer`’s https backend does not wrap itself) both become `PveApiError`. No method caches anything, and the per-run topology cache belongs to `topology.py` (section 3.5), which is also where the bounded thread pool for the  $O(\text{VMs})$  config fetch lives (60-topology.md, REVIEW.md P-02).

`_call()` does retry exactly one failure mode: `proxmoxer.AuthenticationError` gets one reauthenticate-and-retry cycle (REVIEW.md P-01) before becoming a `PveApiError`, via a reauthenticate callback `build_client()` wires in — a full fresh login through `_build_api()` again, not a reuse of the rejected ticket. This is not a general retry policy (a `ResourceException` or a transport failure still fails immediately, on the first attempt); it exists specifically because a ticket can expire for reasons with nothing to do with the call that hits it — a long `apply --confirm` wait, a suspended process, a clock jump — and `proxmoxer`'s own lazy per-request renewal only notices the age of its own clock, not whether the server-side ticket actually outlived a gap that long. A test double built with a bare `PveClient(fake_api)` (no `reauthenticate=`) gets none of this — it behaves exactly as it did before P-01, which is what every existing fake in `tests/unit/fakes.py` relies on.

`build_client()` also applies `proxmox.ticket_refresh_seconds` to the constructed session, best-effort: `proxmoxer 2.x` has no constructor argument for a password/ticket auth's refresh interval (it is a hard-coded `renew_age = 3600` class attribute on `ProxmoxHTTPAuth`, checked lazily on whichever request happens to run after it elapses), so `_apply_ticket_refresh_seconds()` reaches into `api._backend.auth` — undocumented `proxmoxer` internals, not its public interface — to override that instance's `renew_age` after login. Deliberately narrow: it only ever *tightens* `proxmoxer`'s own schedule, never replaces its refresh logic, and if a future `proxmoxer` version reshapes this (or the backend is not `https`, or the auth is API-token, which has no ticket at all), the `getattr/hasattr` guards make it a silent no-op — the P-01 reauthenticate-and-retry above still catches the case this can't reach.

**`storage_definitions()` deliberately uses `GET /storage` (the list form), never the singular `GET /storage/{id}`, even though only one storage's config is needed in some callers.** This was verified empirically against a real PVE 9.2.11 cluster during development ([[dev-cluster-access]] in the maintainer's notes, not reproduced here): an API token granted only `Datastore.Audit` can list every storage's full config — including `saferemove/saferemove_throughput` — via `GET /storage`, but the identical data via `GET /storage/{id}` for the same storage returns `403 Forbidden (Datastore.Allocate)`. The single-item route apparently backs an edit-UI flow gated by the ability to change the config, not merely read it. Since the list form returns the same per-storage object for every storage in one call, there is no reason to ever call the singular form, and using it would quietly force a Storage DRS token to hold more privilege than the tool actually needs. IMPLEMENTATION\_PLAN.md section 3.5 was corrected to match once this was found — see that section's note.

**`storage_content()` returned an empty list — not an error — on every storage, node and content-type filter tried, until the token also held `Datastore.Allocate` on that storage.** This was genuinely surprising and worth calling out precisely because it is a *silent* false negative: `Datastore.Audit` alone gives a clean 200 with `{"data": []}`, which looks exactly like “this storage legitimately has no tracked content” rather than “the caller cannot see it.” It was not guessed at — see IMPLEMENTATION\_PLAN.md section 3.5's note, which records the exact before/after (0 entries with `Audit` only, 37 and 1 entries respectively for two real storages once `Datastore.Allocate` was added) on the same live cluster. **The practical consequence: `pve.py`'s own “keep the token to Audit-only” goal for `storage_definitions()` does not extend to `storage_content()` — a Storage DRS token genuinely needs `Datastore.Allocate` (bundled with `Datastore.Audit` in one role) on every storage it manages, or `topology.py` will silently compute `Usext` (section 5.1.1) as if every storage were empty of untracked volumes, with no error to notice by. docs/manual/00-installation.md states the exact minimum ACL grant an operator needs; this is not a place to economize on privilege out of a preference that turned out not to match reality.**

Two more things about Proxmox's ACL model this surfaced, both now reflected in the manual's token-setup instructions: **an API token's effective permission is the intersection of its owning user's permissions and whatever is granted to the token itself** when privilege separation is on (the default) — granting a role to only one of the two has no effect, both need it; and the grant that was confirmed to work was

made explicitly at each `/storage/{id}` rather than only at the parent `/storage` path, so that is what the manual instructs, rather than relying on propagation that was not cleanly isolated as working or not.

### **`move_disk()`: the one and only bytes/s -> KiB/s conversion**

Section 9.2 is explicit that `bwlimit`'s bytes/s-to-KiB/s conversion happens “at the call site and nowhere else.” `PveClient.move_disk()` is that call site: every caller in this codebase, present and future, passes `bwlimit_bytes_per_sec` in the same unit migration.`bwlimit_bytes_per_sec` already uses, and the division by 1024 (rounded, since the API wants an integer) happens exactly once, inside this method. `format` defaults to `None` and is only ever forwarded when a caller explicitly passes one — `move_disk` without `format=` preserves the source format, which is what section 3.5 requires by default.

## **Building the disk/storage/group model**

**What does this page answer?** How does `topology.py` turn `pve.py`'s raw API responses into the per-group D/S sets the rest of the engine needs, and where does each `IMPLEMENTATION_PLAN.md` section 5.3 (C2) pin condition actually get decided? Describes `proxmox_storage_drs/topology.py` and `proxmox_storage_drs/reserve.py`.

### **D means every disk this module places, pinned or not**

This is stated at the top of `topology.py`'s own docstring because it is the single easiest thing to get backwards: D (section 5.1) is *every* disk this module puts into a group's `Group.disks`, whether or not (C2) pins its placement. A disk this module never sees at all — a stopped VM excluded by `exclude.running_only`, or a disk whose current storage is not in any configured group — is what “foreign” (`Userext`, section 5.1.1) means. Config-excluded disks (`exclude.vmid`s/`exclude.disks`/tags) are *not* foreign: (C2) pins them into D specifically so their bytes still count in (C4)/(C5) and their fragmentation toward  $\kappa$  (section 3.6, “Pinned disks are modelled, not ignored”). An earlier draft of section 5.1.1 listed config-excluded disks as foreign, which directly contradicted (C2) — that self-contradiction was found and fixed in the same commit that first implemented this join; see that section's note if you need the history.

### **One pass, in the order section 3.5 lists**

`build_topology()` fetches everything it needs exactly once, in five stages (`_fetch_cluster_data`, then a loop over `client.vm_resources()` calling `_collect_vm_disks` per VM, then `_build_storages`):

1. `storage_definitions()` (the list form — see `50-pve-api.md`), which is also what every `/.../ storage` pattern (section 11.4, below) is matched against, then the configured groups are expanded and validated against it (`_expand_and_validate_groups`): a literal storage that does not exist, or that lacks images in its content types, is a fatal `TopologyError` (section 11.1: “storage ids exist in the cluster”). A storage that exists but is not marked shared is a warning, not a fatal error — plausible, if unusual, for a single-node group.
2. `storage_resources()` once, to pick one active node per storage (`_pick_active_node`, preferring one reporting status: “available”).
3. `storage_content()` and `storage_status()` per group storage.
4. `vm_config()` and `vm_snapshots()` per considered VM. A VM is *not* fetched at all when `exclude.running_only` is set and it is stopped — the one case `cluster/resources`'s own fields can decide without a config fetch. Every other VM is fetched, including config-excluded ones: whether a specific disk turns out to be “ungrouped” is only knowable after seeing which storage it is actually on, and a config-excluded VM's disk still needs its size for (C4)/(C5) (previous section).

This is the one step that runs across a bounded thread pool (`config.proxmox.read_workers`, REVIEW.md P-02): section 3.5's whole point is that `GET .../qemu/{vmid}/config` has no batch form,

so for a several-hundred-VM cluster this is the read phase concurrency actually helps. It is split into `_fetch_vm()` (network I/O only, run by the pool, one call per considered VM) and `_join_vm_disks()` (pure — the same function the old sequential `_collect_vm_disks()` was renamed from, unchanged in what it computes). `ThreadPoolExecutor.map()` returns each VM's fetch in the same order `client.vm_resources()` listed it, even though the fetches themselves complete in whatever order the pool schedules them, so the join phase runs single-threaded and in the original order — `disks_by_group`, `referenced_volds` and `warnings` end up byte-for-byte identical to a fully sequential run regardless of `read_workers` or of thread scheduling, and never need a lock. `PveClient` itself may reauthenticate mid-fetch if several threads hit an expired ticket at once (50-pve-api.md's P-01 note); that reauthentication is what its own lock serializes, not this join.

5. Non-QEMU resources (`type != "qemu"`) are skipped outright — this tool never touches LXC containers, and section 3.5's read/write paths are qemu-only throughout.

### The lock field comes from `vm_config()`, not a separate call

Section 9.3's own pseudocode says `lock` is readable from either `/qemu/{vmid}/config` or `/status/current`. `_collect_vm_disks` reads it from the config response already being fetched for the disk join, so planning-time lock detection costs no extra API call. `/status/current` is still needed later, but only immediately before a specific move (section 9.2's pre-move re-validation, `execute.py`, not yet written) — never here.

### Pin priority: `_pin_reason()`

One function, checked in the fixed order section 5.3 (C2) lists its conditions: config exclusion (VM-level, then `exclude.disks`), a real snapshot or an unreferenced companion volume (section 3.7, `_disk_snapshot_or_orphan_reason` — the "current" entry `GET .../snapshot` always returns, verified on a live cluster, is filtered out before counting), the per-disk cooldown (below), a VM lock, then an unusedN disk when `exclude.include_unused_disks` is false. A disk gets at most one reason; the first that applies wins, matching how an operator would explain it.

### The per-disk cooldown pin

`build_topology()` takes an optional state: `state.State` and `now: datetime` (both default to "no cooldowns"/the real clock — see `docs/internals/15-state.md`) and computes `state.active_disk_cooldowns()` once per group, before the per-VM join loop, into `cooldowns_by_group: dict[str, dict[str, float]]` (bare `vmid:device` → seconds remaining). Inside the loop, once a disk's `group_name` is known, `cooldowns_by_group[group_name].get(key, 0.0)` is passed to `_pin_reason()` as `cooldown_remaining_seconds` — a plain float, not a second `State` lookup inside that function, keeping `_pin_reason()` a pure decision over already-resolved flags exactly like every other condition it checks. The reported reason names how much of `gates.cooldown_per_disk` remains: `cooldown: moved recently, 1.0h left on gates.cooldown_per_disk`.

This pin is **not** exempted when the disk's storage is actively violating (C4)/(C5) — see `docs/internals/15-state.md`'s "Deliberately not implemented" section for why, and why that gap is bounded and safe rather than silently wrong.

### Sizes: content is authoritative, config is the fallback

`_resolve_disk_size_and_format` looks up the volume in that storage's content listing first (section 3.5: authoritative for what is actually allocated) and falls back to parsing the VM config's own `size=` only when the volume is missing from content — logged as a warning naming the disk, since assume-thick-provisioning sizing from a stale or unlisted config value is a real, if rare, source of drift. Two independently verified reasons this fallback path is not merely defensive: the `Datastore.Allocate-vs-Audit` privilege gap (50-pve-api.md), and an unusedN volume, deleted directly on the storage backend outside Proxmox (confirmed with the cluster's operator), that PVE's own config still referenced and that PVE itself never

noticed or refused to boot the VM over — IMPLEMENTATION\_PLAN.md section 3.6’s note. `_parse_pve_config_size_bytes` parses PVE’s own config-file size suffixes (512G meaning binary GiB, no explicit i) — deliberately not units.py’s parser, which is for *this project’s* config file, a different and coincidentally similarly-shaped format.

### U<sub>s</sub>e<sub>x</sub>t: everything not referenced

Every disk actually placed into a group’s D records its valid in `referenced_volid`s[storage\_id] as it is processed. `_build_storages` then sums the size of every content entry on that storage whose valid was never referenced — orphans, templates, other groups’ foreign volumes, and disks of VMs this run never fetched (stopped-and-excluded, or ungrouped) all fall out of this one computation with no special-casing per category, which is exactly section 5.1.1’s definition read literally.

### Section 11.4: /.../ storage patterns are expanded here, once

`groups[].storages[].id` may be a /.../ pattern instead of a literal id. Expansion happens inside `_fetch_cluster_data`, right after `storage_definitions()` returns, and is the one thing that runs *before* the ordinary literal-id validation described above — `_expand_group()` matches every pattern entry in a group against the sorted list of the cluster’s storage ids with `re.fullmatch` (whole-id, case-sensitive: `/prod/` must not also catch `preprod`), then folds in the group’s literal entries on top, a literal always overwriting whatever a pattern matched for that same id in a plain dict assignment — precedence is a dict overwrite, not a positional rule, so entry order in the file never matters.

Two things a pattern can get wrong are checked right there, before anything downstream sees the result: a pattern matching zero storages is a `TopologyError` (`_match_pattern_entries`, same treatment section 11.1 already gives a literal id that doesn’t exist — a typo is the likelier cause), and two *different* patterns in the same group matching the same storage is also fatal (`claimed_by` in `_match_pattern_entries`) — which entry’s options should apply would be arbitrary. A third check, `_check_cross_group_uniqueness`, runs once every group has been expanded: it is what catches a pattern in one group and a literal (or another pattern) in a different group matching the same real storage — `config._check_group_storage_membership` cannot see that case at load time because it only ever compares two groups’ raw entry *text*, and `"san-a"` and `"/san-.*/"` are different text even when they’d resolve to the same storage.

One deliberate asymmetry: a pattern-matched storage does **not** get the literal path’s “must have images in its content list” check. An operator who names a storage explicitly gets a hard error for a typo; a pattern that happens to pick up an ISO-only or otherwise unusable LUN is a balance-quality concern, not a safety one — (C2) still keeps any disk from ever landing there, and  $u^* = L_s / c_s$  (section 4) simply divides by one more `c_s` than intended. This is called out explicitly in section 11.4’s own closing paragraph, not an oversight here.

Every pattern’s expansion — the raw entry and the ids it matched — is logged at INFO (`logger.info(..., extra={"event": "storage_pattern_expanded", ...})`) and carried on `Topology.pattern_expansions` for `cli.py`’s `verify-storages` to print (section 3.5). `Topology.unmanaged_storage_ids` is computed once, right after groups — every cluster storage id `storage_definitions()` returned minus every id `storage_group_of` resolved to — and is the storage-level twin of the per-disk “ungrouped, not managed” warning `_join_vm_disks` already emits: it does not need a disk to actually be sitting on the storage to be visible.

By the time `_build_storages` runs, `data.expanded_by_group[group.name]` holds real storage ids only — it iterates that, not `group_cfg.storages`, so `Storage.id` and everything keyed off it (`storage_group_of`, `cooldown` keys, `state.json`) never carries pattern text. Nothing past `_expand_and_validate_groups` in this module, or in any module downstream of it, needs to know a pattern was ever involved.

**reserve.py: one (C4)/(C5) evaluator, shared**

`compute_reserve_status()` is deliberately its own module, not a method on `Storage`: the solver (`optimize.py/heuristic.py`, not yet written) will need to evaluate the identical formula against a *candidate* assignment, not only the current one, and `pve-storage-drs show-load`'s reporting needs it against the current assignment today. Both call the same function (AGENTS.md section 5) — `show-load`'s use of it is already exercised end-to-end (see `cli.py`'s `_render_show_load_human/_json`). The section 14 worked example's initial state (`tests/unit/test_reserve.py`) is the proof this reproduces the plan's own arithmetic exactly, not merely a self-consistent unit test.

**The load model**

**What does this page answer?** How does `loadmodel.py` turn six raw Prometheus series into the single per-disk  $\ell_d$  section 4 defines, and what happens when the data for one disk cannot be trusted? Describes `proxmox_storage_drs/loadmodel.py`.

**One function, one group, already-built input**

`compute_group_load()` takes an already-built `topology.Group` (one group's D/S, section 5.1) and a `metrics.PrometheusClient` and returns a `GroupLoad`: one `DiskLoad` per disk ( $\ell_d$ ) and one `StorageLoad` per storage ( $\ell_s, u_s$ ), plus the group's  $u^*$  and whether the whole group is idle. It never touches the PVE API and never builds topology itself — matching `reserve.py`'s identical “given the join, evaluate one formula” framing, not a coincidence: both are meant to be called again, unchanged, by the solver once it exists.

**Six queries, not six-times-groups queries**

The six raw quantities (section 3.4) are fetched with **one instant query each**, unfiltered by group — `sum by (vmid, device) (...)` already returns every disk Prometheus currently reports, and picking out one group's keys in Python is free, whereas six *more* Prometheus queries per group is not. `read_time_ns/write_time_ns` are converted from nanoseconds to seconds here, engine-side (`_combined_raw`), for the same reason F-03 moved `read_factor/write_factor` engine-side: one tested constant beats the same `/1e9` repeated across deployed PromQL strings.

**Known limitation, not a bug:** a config with more than one group re-runs these same six unfiltered queries once per group — correct (each group's coverage/rejection decisions are independent) but wasteful for a many-group cluster. Fetching once and slicing per group would be a straightforward follow-up if a real deployment's group count ever makes this Prometheus load worth avoiding; nothing about `compute_group_load()`'s per-group signature forces the current one-fetch-per-call shape, it is just what the first implementation does.

**Coverage rejection excludes a disk from the group total, not just from its own  $\ell_d$** 

Section 3.4: “reject a disk whose sample coverage over  $W$  is below `window.min_coverage` and fall back to its last known load from `state.json`... never treat missing data as zero load.” Two things worth being explicit about, since the plan states the rule but not this mechanism:

1. **The rejected disk's raw contribution is excluded from  $T_g/0_g/B_g$  entirely**, not merely capped or zeroed in place — one noisy or half-missing series must not bias every *other* disk's normalized share.  $\ell_d$  for every accepted disk is computed from the accepted-only totals.
2. **The rejected disk's own  $\ell_d$  is substituted afterward**, from `last_known_loads` if the caller has one (keyed by `topology.Disk.key`), or flagged `0.0` if not. `DiskLoad.flagged_reason` is set either way, so a caller can never mistake a flagged `0.0` for a genuinely idle disk — `show-load` renders it as an explicit  $\Delta$  line, never silently.

state.py (section 11.2) now provides the real last\_known\_loads source: cli.py’s show-load and plan both pass state.load\_vector\_for\_group(state, group.name) through unchanged, exactly the mapping this function already expected (see docs/internals/15-state.md). compute\_group\_load() itself needed no change at all — the parameter was shaped for this from the start. A group with no recorded balance yet (a fresh state.json, or none on disk) still gets None, and every coverage-rejected disk with no seed is flagged 0.0, exactly as before.

### **tpmstate0/unusedN are never rejected; efidisk0 is**

Section 3.4’s own note: tpmstate0 and unusedN are not QEMU block devices and legitimately emit no blockstat series — `l_d = 0` for them is correct, not a data-quality problem. `_is_metrics_expected_absent()` encodes exactly this device-name distinction so these two kinds are never flagged for low coverage, while efidisk0 — a real QEMU drive that *should* appear — is held to the same `min_coverage` bar as `scsi/virtio/ide/sata`. Getting this backwards either way is a real bug: exempting efidisk0 would hide a genuine metrics gap on it; not exempting tpmstate0/unusedN would spam a flag on every run for every cluster, for a “gap” that isn’t one.

### **T\_g = 0 is “idle”, computed only from accepted disks**

GroupLoad.idle is True exactly when every *accepted* disk’s `raw_t` is zero — the guard section 4 requires (“if `T_g = 0` the entire group is idle, skip it”). An all-rejected group (e.g. a total Prometheus outage mid-window) also comes out `idle=True` by this same computation, which is not quite literally accurate (the truth is “unknown”, not “idle”) but is harmless: both mean “do not act”, which is the only decision that follows from either. A caller that wants to distinguish the two can — every disk’s `flagged_reason` is still there to check.

### **L\_s/u\_s are the *current* assignment, not a candidate one**

StorageLoad sums `l_d` over disks whose `Disk.current_storage` already points at that storage — today’s state, exactly like `reserve.py`’s `compute_reserve_status()`. The solver (`optimize.py/heuristic.py`, not yet written) will evaluate the identical per-disk `l_d` values against a *candidate* assignment instead; nothing in `loadmodel.py` needs to change for that, since `l_d` itself does not depend on which storage a disk is currently on.

### **compute\_disk\_load\_series(): the same blend, as a time series**

`compute_group_load()` reduces each raw quantity to one already -quantile’d scalar per disk via `quantile_over_time` at query time. Section 10’s forecaster needs the opposite: the *raw* per-disk `l_d` signal, unreduced, over a range and step of its own choosing (typically `forecast.required_range_seconds()`, not `window.lookback` — section 10.1 is explicit these are “genuinely different things”). `compute_disk_load_series()` fetches the identical six `rate(...)` expressions `_fetch_raw_quantity()` builds, `query_range`’d instead of wrapped in `quantile_over_time` and `instant_query`’d, then runs section 4’s exact `normalize-then-weight-then-rescale` blend once **per timestamp** instead of once. `_combine_raw_values()/_blend_loads()` are that formula factored out to plain-number arguments so both this function and `compute_group_load()` call the identical implementation (AGENTS.md section 5) — extracted, then proven behavior-preserving by `test_loadmodel.py`’s full existing suite passing unchanged before any series-specific test was added, the same regression-by-refactor discipline `reserve.transient_charge_ok()`’s own extraction used.

Every disk in the group gets an entry, even an empty one — unlike `compute_group_load()`, this function does not apply `window.min_coverage` at all: a forecaster’s own `required_range()` is a much longer, coarser signal than that rule was built to validate, and a sparse history is exactly what the forecaster itself needs to see to distrust its own fit. Not called from anywhere yet — this is section 10’s raw material for the

section 7.3 saturation guard `payback.py` does not implement yet either (see that page’s own note on the gap); this function exists so that work has its data source in place first.

### What is still missing from phase 3

IMPLEMENTATION\_PLAN.md section 12 phase 3’s own “done when” is “correct act/no-act decision per group, with the reasoning shown” — that is the section 6 drift/imbalance gates, cooldowns, and the reserve override, none of which exist yet. `loadmodel.py` is the load computation those gates consume, not the gates themselves; `pve-storage-drs show-load` reports `l_d/L_s/u_s` today, but no command yet decides whether a group should be re-balanced.

### Gating: deciding whether to act at all

**What does this page answer?** How does `gates.py` turn a `loadmodel.GroupLoad` and per-storage `reserve.ReserveStatus` into the section 6 act/no-act verdict, and why isn’t cooldown handling here? Describes `proxmox_storage_drs/gates.py`.

### One function, three gates, in the order section 6 lists them

`evaluate_group_gates()` is pure — no network I/O, no `state.json` access — taking a group’s already-computed `GroupLoad`, a `storage_id` → `ReserveStatus` mapping (the caller’s job, via `reserve.compute_reserve_status()` — this module never recomputes it, per AGENTS.md section 5’s “one implementation of every rule”), the group’s `GatesConfig`, and the load vector recorded at the group’s last *executed* balance (`last_load`, `None` if there isn’t one). It checks, in section 6’s own order, stopping at the first that decides:

1. **Reserve override.** Any storage in `reserve_statuses` with `.violated` set forces `act=True` immediately, bypassing drift and imbalance entirely — section 13’s “safety is not subject to hysteresis” made literal, not just documented intent.
2. **Drift gate.**  $\|l_{\text{now}} - l_{\text{last}}\|_1 / \|l_{\text{last}}\|_1 \geq \text{gates.drift\_threshold}$ , vectors aligned over the **union** of disk keys (a disk absent from one side contributes its full load in the other as drift — a new or deleted disk is a genuine change to the group’s I/O profile). `last_load=None` skips this gate outright, per section 6’s own degenerate-case table — not “treat as zero drift”, which would make an operator’s very first run fail to act on an already-imbalanced cluster.
3. **Imbalance gate.**  $(\max_s u_s - \min_s u_s) / u^* \geq \text{gates.imbalance\_threshold}$ ,  $u^*$  from `GroupLoad.average_utilization`.  $u^* == 0$  (idle group, section 4) always means no-act — there is nothing to spread evenly.

### last\_load is real now, for show-load and plan

`state.py` (section 11.2) exists: `cli.py`’s `_last_loads_by_group()` reads `state.path` once per run and hands each group’s slice of `last_balance.load_vector` (re-keyed to plain `topology.Disk` key, or `None` if that group has no recorded balance yet) to `evaluate_group_gates()` as `last_load` — both `show-load` and `plan` do this identically (`cli.py` calls the one function, not two copies). `evaluate_group_gates()` itself needed no change at all: it already took `last_load` as a parameter and treated `None` correctly (see `docs/internals/15-state.md` for the read side, and this page’s next section for what is still missing).

**What this does not yet do.** `state.json`’s `last_balance` is only ever *read* here — nothing in this codebase calls `state.with_recorded_balance()` yet, because nothing executes a migration yet (`execute.py`, phase 7). Until then, `last_balance.load_vector` for any group stays exactly what it was the last time a human (or a test) wrote it by hand; a config with no `state.json` on disk, or a fresh install, still behaves exactly as this page originally described — `last_load=None`, drift gate skipped, decision reduces to “reserve override, else imbalance”. The wiring is real; the writer that would make it self-sustaining is not built yet.

### show-load's gate line is diagnostic, not a real decision

pve-storage-drs show-load prints each group's verdict (Group <name> → ACT: <reason> / → NO ACTION: <reason>) reusing GateDecision.reason verbatim — deliberately the single source of truth for the wording, rather than a second, similarly-but-not-identically phrased string living in cli.py. Its verdict now reflects real drift history whenever state.json has one, exactly like plan's does — but show-load still isn't what apply (not yet written) will actually gate an execution on; it is plan/apply's own gate evaluation, run at the moment a plan is built or applied, that is authoritative.

### Cooldowns are not this module's job

gates.cooldown\_per\_disk/cooldown\_per\_storage decide which *individual* disks or storages are movable once a plan is already being built — the same kind of decision (C2)'s other pin reasons make in topology.py, not a group-wide act/no-act verdict, so this module still has no cooldown logic of its own. Both are now implemented, in the modules this page always said they belonged with: the per-disk cooldown is a topology.py (C2) pin (docs/internals/60-topology.md), and the per-storage cooldown is a heuristic.py target-eligibility filter (docs/internals/90-heuristic.md), both reading state.py's cooldown data (docs/internals/15-state.md).

### The heuristic solver

**What does this page answer?** How does heuristic.py turn a group's loads into a target assignment, why is the section 5.4 objective its own function rather than baked into the search, and what does this module deliberately not do yet? Describes proxmox\_storage\_drs/heuristic.py.

### evaluate\_assignment() is the objective, kept separate from the search

Section 5.5 requires it explicitly: “the heuristic must use the **same** feasibility and objective functions as the MILP path so the two backends are directly comparable.” evaluate\_assignment() takes a Group and an arbitrary candidate Assignment (dict[Disk.key, storage id], not necessarily reachable from the current state by any sequence of moves) and returns an ObjectiveBreakdown — the section 5.4 objective's five terms kept separate, not collapsed into a single number, because explain (not yet written) needs the arithmetic shown, and because the tests cross-checking this against section 14's worked example need to assert each term individually, not just a total that could be right for the wrong reasons.

This is the same design choice reserve.py made first: a pure function of *data*, with no notion of “the current run” baked in. heuristic.py reuses reserve.compute\_reserve\_status() directly for (C4)/(C5), passing it a storage\_of callback closed over the candidate assignment being evaluated — the extension point reserve.py's own docstring already anticipated (“the solver will pass a candidate assignment through the identical shape”) before this module existed to use it.

### Repair is unconditional, not weight-driven

Section 13: “the reserve is never traded against balance.” The heuristic makes this literal rather than merely well-weighted: \_repair() runs *before* \_descend() and fixes any (C5) violation by moving disks off the worst-violating storage, one iteration at a time, regardless of what the objective weights say — there is no beta/gamma value large enough to switch this off. objective.reserve\_violation\_penalty still appears in ObjectiveBreakdown.reserve\_penalty\_term, but only for *reporting* a residual that repair could not fix (physically impossible, not merely unattractive) — by the time \_descend() runs, every reserve violation \_repair() could resolve already has been, so in the common case this term is exactly zero and does not influence descend's choices at all. test\_reserve\_violation\_is\_repaired\_even\_with\_beta\_high\_enough\_to\_forbid\_balance\_moves is the regression test for this: beta=100 (so no purely-cosmetic balance move could ever be worth a migration) still sees the repair move happen.

**A candidate is judged by the group-wide total shortfall, not the source storage's own.** Moving a disk off a violating storage always reduces *that* storage's own shortfall — it has fewer bytes, and if anything a smaller largest-disk requirement — which makes “does the source's shortfall fall” a tempting but wrong test: it says yes even when the disk only relocates the problem to whichever storage receives it, or makes the group's total worse. An earlier version of `_repair` used exactly that test and would oscillate a disk back and forth between two storages neither of which can fully hold it, “successfully” reducing the current worst offender's shortfall on every iteration while never making net progress. Requiring the *group* total to strictly decrease fixes this and is also what makes the loop's iteration bound (movable disks times storages, generously sized rather than tightly) actually a correct termination argument rather than a lucky one. `test_repair_does_not_oscillate_when_no_target_can_fully_absorb_the_violation` is the regression test: a 3 TiB disk that cannot fit, alone, on either of two 8 TiB/`reserve_factor`: 2.0 storages is moved exactly once (repairing what can be repaired, from a 3 TiB group-wide shortfall down to 1 TiB), not shuffled back and forth forever.

### Descend explores swaps, not only single moves

Section 5.5 is explicit that swaps are not an optional refinement: “when both storages are near their capacity limit, no single move is feasible, and only an exchange of two disks can improve the balance.” `_descend()` evaluates every single-disk move *and* every pairwise swap of two movable disks each iteration, applying whichever most reduces the objective, and stopping when nothing does or after `heuristic_` iterations. Both are plain re-evaluations of `evaluate_assignment()` against a trial dictionary — no separate swap-feasibility logic exists, because a swapped assignment is just another candidate the same objective function scores. `test_descend_uses_a_swap_when_no_single_move_is_feasible` constructs exactly the capacity-locked scenario the plan describes (two storages each with headroom smaller than any single disk) and checks the only reachable improvement — a swap — is the one found.

### Descend also tries relocating a whole VM at once — a real gap, found by dogfooding

Neither of the two candidates above can reach a third kind of improving move: relocating *every* one of a multi-disk VM's movable disks to one new storage, together. A single-disk move can't get there when the objective's  $\kappa$  (VM affinity) term is large enough to make moving just one of the VM's disks a net loss on its own — it temporarily *fragments* the VM (paying  $\kappa$ ) before a second move could reunite it elsewhere, and `_descend()` only ever takes a step that is itself improving. A swap can't either: it exchanges two disks' positions with *each other*, never relocates a whole disk set to a third storage together.

This was not a theoretical gap — it was found by testing against a real production cluster (not by review): a two-disk VM sat entirely on one storage that also carried a large amount of *pinned* (immovable) load, with `kappa_vm_affinity` at its default 0.50. Every individual disk move `_descend()` tried was a net loss once fragmentation was paid for, so it found *zero* improving moves at all and left the cluster at its full, 196%-imbalance starting point — even though CP-SAT, given the exact same data, found the two-disk relocation immediately. A dependency-free heuristic that can get stuck this badly on a real, unremarkable cluster undermines the one thing section 5.5 asks it to be: “the path for very large groups,” reliable with no optional solver installed at all.

The fix adds a third candidate family, tried in the same “evaluate everything, apply whichever wins” step as the other two: `_vm_relocation_candidates()` groups movable disks by `vmid` (skipping any VM with only one movable disk — that case is already the plain single-move candidate), and for each such VM, `_descend()` tries moving *all* of its disks to each other storage in one trial. Scored by the exact same `evaluate_assignment()` call as everything else, so it can never itself choose a worse assignment, and it respects `cooldown_storages` as a destination exclusion the same way the other two candidates do. `test_descend_relocates_a_whole_multi_disk_vm_neither_single_moves_nor_swaps_can_reach` reproduces the production scenario's numbers by hand (a single move raises the objective by

+0.43; the joint move lowers it by -0.14) and confirms `_descend()` only finds the improvement once this candidate exists. `_best_of()/_single_move_trials()/_swap_trials()/_vm_relocation_trials()` are the flake8-complexity-driven extraction that let all three candidate families share one “keep whichever trial scores lowest” loop rather than three copies of it.

### The storage cooldown excludes a destination, never a source

Section 6: “a storage involved in a migration within `cooldown_per_storage` accepts no new incoming moves.” `run_heuristic()`’s `cooldown_storages` parameter — a plain `frozenset[str]` its caller (`cli.py`, via `state.active_storage_cooldowns()`) computes, not a `state.State` this module reads itself — is threaded through to `_descend()` only. Both of `_descend()`’s neighbourhoods respect it: the single-move loop skips any `target.id` in `cooldown_storages`, and the swap loop skips a pair whenever *either* storage a disk would land on is in cooldown (a swap always sends one disk to each of the two storages it touches, so either side landing on a cooldown storage blocks the whole swap). A disk already resident on a cooldown storage is always free to move *away* from it — the rule is about new arrivals, not existing occupants — which is why the filter only ever appears on the destination side of a candidate, never the source.

`_repair()` never receives `cooldown_storages` at all — a deliberate exemption, not an oversight, matching this module’s own repair-is -unconditional stance above: a storage actively needed to resolve a live (C4)/(C5) violation is not deferred for having been written to recently, the same reserve-override principle `gates.py` and `payback.py` already apply to the drift/imbalance gates and the payback test respectively. `test_run_heuristic_repair_ignores_storage_cooldown` is the regression test: a violation whose only viable full repair target is in `cooldown_storages` is still repaired. `test_descend_blocks_new_arrivals_onto_a_cooldown_storage` and `test_descend_still_allows_a_disk_to_move_away_from_a_cooldown_storage` cover `_descend()`’s own two sides of the rule, both against the section 14 fixture.

### objective.spread\_metric: two different quantities, not one rescaled

Section 5.4/(C6) offers two imbalance forms, and `evaluate_assignment()` implements both: “l1” (default) is  $\alpha * \sum(e_s)$ , every storage’s absolute deviation from  $u * \text{summed}$ ; “minmax” is  $\alpha * \max(u_s)$  — the plan’s own  $t \geq u_s$  for all  $s$ , the raw utilization of the single hottest storage. These are not the same quantity with a different aggregation:  $\max(u_s)$  and  $\max(e_s)$  can disagree, because a storage sitting far below  $u * \text{summed}$  produces a large deviation  $e_s$  that minmax was never meant to notice (the manual’s own words: “indifferent to a second nearly-as-bad storage” — nearly-as-bad on the *high* side; a cold storage is not what minmax was designed to react to at all). `ObjectiveBreakdown` carries both `spread_e` (deviations) and `utilization` (raw  $u_s$ ) regardless of which metric is active, computed once in the same loop, so switching metrics never means computing something the other mode didn’t already have on hand.

REVIEW.md’s Q-01 found this config knob silently ignored by an earlier version of this module (always L1, `objective.spread_metric` never read) — the same class of gap P-01/P-02 found in `pve.py/topology.py`. `test_spread_metric_minmax_is_indifferent_to_a_second_nearly_as_bad_storage` is the regression test, built directly from the manual’s own claim: two scenarios with an identical hottest storage but a materially different second-worst one score identically under minmax and differently under l1.

### Proof this reproduces the plan, not just itself

`tests/unit/test_heuristic.py` doesn’t stop at “the assignment matches the prose.” At the default weights it asserts `ObjectiveBreakdown.total` equals 2.533333 (three-move plan) and, at `beta_move_count: 0.50`, 3.158333 (two-move plan) — the exact totals REVIEW.md Appendix A independently re-derived by hand from the plan’s own numbers, not values this module invented and then asserted against itself. Getting both to five decimal places is strong evidence the objective’s five terms, their units (TiB for size, average in-flight I/O for load), and the search that picks among them are all correct together, not merely internally consistent.

## What this pass deliberately does not do

- **“Polish” (heuristic step 4)** — reuniting a fragmented VM when doing so does not worsen imbalance beyond `gates.imbalance_threshold`. Not implemented: the section 14 fixture’s exact three-move and two-move solutions are both reachable by repair+descend alone (proven by the tests above), because the objective’s own kappa term already makes descend prefer co-location whenever it is not too costly, and (since the fix above) descend’s own whole-VM co-relocation candidate now reaches the two-or-fewer-disk case directly. What’s left for polish is strictly narrower than it once was: an *N-way rotation* across three or more storages (disk A needs  $S1 \rightarrow S2$ , disk B needs  $S2 \rightarrow S3$ , disk C needs  $S3 \rightarrow S1$  in a cycle) where no disk’s own move improves alone and no two disks share a VM — outside every candidate `_descend()` tries today. The one fixture that exists to validate this module does not exercise it. A real gap, tracked here rather than silently absent.
- **(C2) format-compatibility eligibility** — `topology.Storage` does not yet carry the storage type/format information that rule needs (it is resolved internally in `topology.py`’s `_default_format` but never exposed on `Storage`), so every group storage is currently an eligible target for every movable disk. The section 14 fixture is homogeneous (three identically-typed storages) and does not exercise this gap either.
- **CP-SAT/CBC coefficient scaling (section 5.5)** — belongs to `optimize.py`, not yet written. `evaluate_assignment()`’s plain floating-point objective is what the MILP path will need to agree with, not a stand-in for its own separately-scaled integer objective. `heuristic.py` is wired into `cli.py`’s `plan` command, together with `schedule.py` (section 8’s move ordering) — see `95-schedule.md` and `docs/manual/27-plan.md`. `explain/apply` remain stubs (`docs/manual/30-safety-and-status.md`).

## The MILP backends: CP-SAT and CBC

**What does this page answer?** How does `optimize.py` turn section 5.3’s constraint set into a real CP-SAT/CBC model, why is the reserve solved in two lexicographic stages rather than one big-M objective, and what does `cli.py`’s backend dispatch actually do when a MILP solver is unavailable or fails? Describes `proxmox_storage_drs/optimize.py`.

### Both backends solve, then hand off to `heuristic.evaluate_assignment()`

`solve()` builds a model, extracts the winning `x_{d,s}` assignment, and immediately calls the *same* `heuristic.evaluate_assignment()` the heuristic backend uses to produce the reported `ObjectiveBreakdown` — one implementation of the section 5.4 objective, shared by every backend (AGENTS.md section 5; `heuristic.py`’s own module docstring was written anticipating exactly this: “built specifically so `optimize.py` can call the identical function”). A modeling mistake in this module can therefore make the solver choose a *suboptimal* assignment at worst — it can never make the tool *believe* an assignment is better than it is, since the number that ends up in `plan`’s output is never this module’s own objective value.

## Lexicographic, not big-M — and why that is the only mode implemented

Section 5.3 describes two ways to keep a reserve violation from being “traded away” by a large enough balance gain: a single-stage objective with a calibrated penalty  $P$  (big-M), or a two-stage **lexicographic** solve — minimize  $\sum_s r_s$  alone first, then fix that total as a hard constraint and minimize the real objective. `optimize.py` implements only the second, both because the plan itself says to (“use the lexicographic solve by default... needs no calibration, its correctness does not depend on  $T_g$ ”) and because there is no `solver.*` config knob that ever selects big-M at runtime — it exists in the plan text to be *compared against* the lexicographic solve, a comparison tests/fixtures/generate\_expected.py already does by exhaustive enumeration on `reserve-tradeoff.yaml`, not something this module needs a second code

path for.  $\sum r_s > 0$  in a lexicographic result provably means *no* assignment could do better, at any weight — never merely “unattractive”.

**That proof depends on stage 1 actually being solved to proven optimality, not merely to `solver.mip_gap`** (REVIEW.md S-07). An earlier revision applied the same `mip_gap` to both stages: CP-SAT’s stage 1 accepted FEASIBLE (an incumbent, not a proven minimum) whenever the gap was satisfied, and PuLP’s own `LpStatus` string is “Optimal” whether CBC proved the bound or merely stopped because `gapRel` was satisfied — either way, stage 2 then pinned  $\sum r_s$  to that unproven incumbent (`==`), which can both accept a shortfall up to the gap *and* forbid finding less of it. Fixed by forcing stage 1’s own gap to 0 in both backends regardless of the configured `solver.mip_gap` (which now applies to stage 2’s real objective alone), requiring CP-SAT’s stage 1 status to be exactly OPTIMAL (never FEASIBLE), and changing stage 2’s slack constraint from `==` to `<=` in both backends — monotone-safe (stage 2 can never do worse than the now-proven stage-1 value) rather than brittle against a hypothetical disagreement between the two solves. Cheap in practice: stage 1 is a near-feasibility problem that closes instantly on realistic groups, so forcing an exact proof costs nothing measurable.

### CP-SAT: section 5.5’s exact integer scaling

Every coefficient is folded and rounded exactly as section 5.5 specifies:

- `_LOAD_SCALE` ( $K = 10^6$ ) scales every load-valued quantity ( $e_s, t, u^*$ ). (C6)’s per-(disk, storage) coefficient is  $a_{\{d,s\}} = \text{round}(K \cdot \ell_d / c_s)$  — folded and rounded *once*, per the plan’s own warning against computing  $\text{round}(K/c_s)$  and multiplying separately (which rounds the capability weight itself and would make CP-SAT and CBC disagree for a non-integer  $c_s$ ).
- `_WEIGHT_SCALE` ( $W = 10^4$ ) scales every objective weight ( $\alpha/\beta/\gamma/\kappa$ ). The  $\gamma$  term folds  $z_d$  into its own coefficient ( $\text{round}(\gamma \cdot W \cdot K \cdot z_d^{\text{TiB}})$ ) rather than factoring out a standalone  $\gamma_{\text{scaled}} = \text{round}(\gamma \cdot K / 2^{20})$  — the plan’s own worked example of *why* that shortcut is wrong: at the default  $\gamma = 0.05/\text{TiB}$ , it rounds to zero and CP-SAT would silently stop caring about disk size when choosing what to move.
- Every size-valued quantity ( $Z_s, R_s, r_s, z_d, C_s, U^{\text{sect}}, \text{min\_free\_bytes}$ ) is a whole-MiB integer (`_mib()`), already exact, needing no scale of its own.
- `_RESERVE_FACTOR_SCALE` is this module’s own addition, not named in the plan: (C5)’s  $R_s \geq f_s \cdot Z_s$  multiplies a *variable* ( $Z_s$ ) by `reserve_factor`, which section 5.5’s “fold constants into a coefficient” recipe does not directly cover (that recipe is for a constant multiplying a *binary* decision variable). Expressed instead as  $\text{SCALE} \cdot R_s \geq \text{round}(f_s \cdot \text{SCALE}) \cdot Z_s$  with  $\text{SCALE} = 10^6$  — a single linear constraint, correct to within  $1/\text{SCALE}$  MiB, utterly negligible next to any real disk size.

Section 5.5 also asks for two build-time assertions guarding this folding, and `_cpsat_objective_terms()` now has both (REVIEW.md S-09): `_assert_nonzero_when_weighted()` catches exactly the  $\gamma$  trap the paragraph above describes — a non-zero configured weight whose *rounded* coefficient collapsed to 0 (only ever a sub-kilobyte disk at the default  $\gamma$ , per `test_gamma_trap_assertion_fires_end_to_end_for_a_sub_kilobyte_disk`) — and `_assert_objective_magnitude_within_int64()` checks a coarse, deliberately conservative worst-case sum of every term against  $2^{62}$ , an order of magnitude under CP-SAT’s own  $2^{63}-1$  domain limit. Neither applies to CBC, whose continuous, unscaled model has no analogous overflow risk.

`AddHint(x[d, o(d)], 1)` warm-starts every candidate from the current assignment, in both stages (harmless when unused, and stage 1’s own optimum is frequently “keep everyone where they are” when nothing already violates (C5)).

### CBC: “direct transcription”, deliberately not scaled

The plan’s own words for this backend: “continuous  $e_s, Z_s, r_s$  are fine.” `_cbc_feasibility_constraints()/_cbc_objective_terms()` write the identical constraints CP-SAT does, but with

plain floats — no `_LOAD_SCALE/_WEIGHT_SCALE` anywhere, because CBC needs none of CP-SAT’s integer-coefficient discipline. `solver.mip_gap` and `solver.time_limit_seconds` map onto `pulp.COIN_CMD(gapRel=..., timeLimit=...)` directly.

Every variable is built with `_lp_variable()`, a thin wrapper around the direct `pulp.LpVariable(name, ...)` constructor — **not** `prob.add_variable(...)`, PuLP v4’s replacement, despite an earlier version of this module having switched to it (REVIEW.md S-01, found by running the CBC tests against the exact pulp version Debian trixie packages). `LpProblem.add_variable` does not exist before pulp 3.3.1; trixie packages 2.7.0, and section 2.1 designates CBC-through-python3-pulp as *the* packaged solver path, so the “modernize to v4” choice crashed on the platform this project targets, while working fine — and passing every test — against a hand-picked newer pulp installed by hand. Direct `LpVariable(...)` construction works unchanged across every pulp version from 2.7.0 through 3.3.2 (bisected by wheel inspection); on 3.3+ it also raises a `DeprecationWarning` recommending `add_variable`, which `_lp_variable()` suppresses explicitly with a `warnings.catch_warnings()` block — the warning is real, but the alternative it recommends is the one that cannot run on the deployment target, so silencing it is the correct fix, not a workaround. `_pulp_solve()` similarly turns a `PuLPSolverError` (Debian splits the `coinor-cbc` binary into a separate Recommends, not a Depends, of `python3-pulp` — the library can import with no working solver behind it) into this module’s usual `None`, matching `solve()`’s “never raises, `cli.py`’s fallback cascade handles it” contract instead of a bare traceback.

### (C1)/(C3)/(C4)/(C5) are shared code, factored once per backend

`_cpsat_feasibility_constraints()/_cbc_feasibility_constraints()` build exactly the constraints both lexicographic stages need identically — only the *objective* differs between stage 1 (Minimize( $\sum r_s$ )) and stage 2 (the real section 5.4 objective, with  $\sum r_s$  fixed to stage 1’s result as a hard constraint). Each is called twice, with fresh variables each time, rather than trying to mutate one model’s objective in place — simpler to reason about than partial reuse, at the cost of building the constraint set twice per solve (irrelevant next to a MILP solve’s own running time).

(C3)’s `y_{v,s}` linking honors `objective.affinity_counts_pinned_disks` exactly as the plan’s own text describes it: **false** (default) links `y` to `x` over  $D^{\text{mov}}$  only, so a pinned disk never forces `y_{v,s}=1` for its storage; **true** ranges the constraint over all of  $D$ , which this module implements as forcing `y_{v,s}=1` outright wherever a pinned disk of `v` already sits (the natural reading of “ $x_{\{d,s\}} \leq y_{v,s}$  for  $d \in D$ ” when `x_{\{d,s\}}` is a constant 1, not a variable, for a pinned disk).

### cli.py’s dispatch: auto cascades, an explicit backend still falls back

`cli._solve_group()` implements `solver.backend`’s four values: “auto” tries `cpsat` then `cbc`; “`cpsat`”/“`cbc`” try only that one; “heuristic” skips `optimize.solve()` entirely. Whichever backend `solve()` cannot use (library not importable, or no feasible solution within `solver.time_limit_seconds`) returns `None` — never an exception — and `_solve_group()` moves to the next one in the cascade, ending at `heuristic.run_heuristic()` if every MILP attempt failed. This applies **even to an explicitly forced backend**: section 13’s failure-mode table says “solver infeasible or timing out -> fall back to the heuristic; never emit a partial/unvalidated assignment”, with no carve-out for `solver.backend: cpsat` specifically. The one difference: falling back from an *explicit* backend logs a warning (an operator who named `cpsat` to reproduce a specific result should not have to diff --json output to notice auto quietly ran instead); falling back within auto itself is expected and silent. plan’s human and --json output both report which backend actually produced each group’s plan (`solver: cpsat (optimal) / "solver_backend": "cpsat", "solver_status": "optimal"`), precisely so that distinction is never hidden.

### The storage cooldown: enforced in both stages, both backends (REVIEW.md S-04)

`_cpsat_feasibility_constraints()/_cbc_feasibility_constraints()` both take `cooldown_storages` now and add one constraint per  $(d, s)$  pair where `s` is in it and is not `d`’s current storage: `x_{\{d,s\}} = 0`.

Applied inside the shared feasibility-constraint builder, it lands in *both* lexicographic stages automatically (built once per stage, per `solve()`'s own architecture) — a disk cannot be assigned to a cooldown storage it is not already on, in either the reserve-minimization stage or the real-objective one, matching section 6's "accepts no new incoming moves" and `heuristic._descend()`'s own identical `target.id` in `cooldown_storages` check. A disk already resident on a cooldown storage is left free to leave it, or to stay — the cooldown only ever blocks an *arrival*.

An earlier revision of this module left the parameter accepted but unenforced, reasoning that cooldowns were inert anyway since nothing wrote `state.json`'s cooldown timestamps yet. That reasoning stopped holding the moment `execute.py/apply` (phase 7) started calling `state.with_recorded_cooldown()` — from that point on, `solver.backend: auto`'s default cascade to CP-SAT/CBC could plan straight through a cooldown the heuristic would have respected, silently disagreeing with it on the exact case the cooldown exists for. Fixed as above.

**Not replicated:** `heuristic._repair()`'s own exemption from this same cooldown when fixing a live (C4)/(C5) violation (section 13's reserve-override principle — a repair move is never deferred for a cooldown). The fix here is one constraint applied uniformly to both solve stages, including stage 1's reserve-shortfall minimization itself, so a cooldown storage stays excluded even when it is the *only* way to resolve a violation — narrower than the heuristic's behaviour in that one specific edge case (a group simultaneously mid-violation and mid-cooldown on its one viable target). Reported honestly rather than silently: stage 1's own slack is genuinely 0 achievable-within-constraints with the cooldown storage excluded, so plan still reports an accurate number, just a more conservative one than the heuristic would compute for the identical input. Replicating the exemption exactly would need a second, repair-only relaxation of the same constraint or a restructured two-phase solve — a real, separately-scoped piece of work.

### Deliberately not implemented in this pass

- **(C2) format-compatibility eligibility** — the same gap `heuristic.py` already documents (`topology.Storage` does not expose storage type/format).
- **A live cluster's numpy/pandas footprint.** `ortools` pulls in `numpy` (and `numpy` pulls in nothing further this project cares about) purely as its own transitive dependency; nothing in this project imports `numpy` directly. A recent `numpy`'s bundled type stubs use syntax newer than this project's `python_version = "3.11"` mypy target, which mypy's `follow_imports = "skip"` (used for `ortools.*` itself) cannot rescue, because a stub-level *parse* error happens before mypy ever gets to apply a per-module setting to it. This surfaces only if `ortools` (via `pip install -e .[solver]`) or `statsmodels`'s own `[forecast]` extra happens to pull in an affected `numpy` version into the *same* venv `make typecheck` runs against — `make install`'s own `.[dev]` never does, so it does not affect this project's actual gate. Both MILP backends were verified for real during development (`pip install -e .[solver]`, all of `test_optimize.py` passing and reproducing the section 14 and `reserve-tradeoff.yaml` fixtures' exact numbers for both CP-SAT and CBC) — see `tests/unit/test_optimize.py`'s own docstring. That first pass installed `pulp 3.3+` from PyPI, not the version actually packaged for the deployment target — a gap `REVIEW.md`'s tenth pass caught (S-01): `python3-pulp 2.7.0+dfsg` (Debian trixie, also what CI's `debian:trixie` job installs) predates `LpProblem.add_variable`, PuLP v4's variable-construction API, which an earlier revision of this module had adopted and which crashed immediately on that version. Fixed by constructing every variable with `pulp.LpVariable(...)` directly (works unchanged on every `pulp` release from 2.7.0 through 3.3.2, bisected by wheel inspection) and suppressing the resulting v4-migration `DeprecationWarning` on 3.3+ explicitly, plus catching `PulpSolverError` around `prob.solve()` (the `coinor-cbc` binary is a separate Debian Recommends, not a Depends, of `python3-pulp` — the library can import with no working solver behind it) so a missing solver returns `None` like every other "this backend cannot produce a plan" case, rather than a bare traceback. `test_optimize.py`'s CBC cases now pass against `pulp`

2.7.0 too, verified directly against a throwaway venv pinned to that exact version — the packaged path, not just the newer one pip happens to resolve.

## Executing a plan: `execute.py` and its `cli.py` wiring

**What does this page answer?** Why does a “completed” move need three separate conditions, not one; why does a VM lock get waited out instead of failing the run; how is the injectable `Clock` used to test all of this without a real sleep; and what exactly does `cli.py`’s apply handler add on top of `execute_plan()` itself? Describes `proxmox_storage_drs/execute.py` and the `_handle_apply()/_plan_group()` half of `proxmox_storage_drs/cli.py`.

### `execute_plan()` re-validates every move against the live cluster, not the plan

`schedule.ScheduleResult.order` was computed against a snapshot of the cluster taken at planning time. By the time a move near the end of a long `confirm` run is about to be issued, that snapshot can be stale in ways a live cluster routinely produces: an operator touched the VM, PVE’s own Dynamic Load Balancer moved it to another node, a snapshot appeared, or another process changed the target storage’s free space. `_preflight()` re-fetches the VM’s current node, config, running status, snapshot list and exclusion tags fresh for every single move — deliberately bypassing the per-run topology cache (`topology.py`’s section 3.5 cache exists to avoid re-fetching for *planning*, not to avoid re-validating immediately before a mutation). The exclusion re-check (section 9.2 step 3: “untagged for exclusion”, REVIEW.md S-06) is `_is_excluded_by_tag_or_vmid()`, a small duplicate of `topology.py`’s own `vmid/tag` predicate rather than an import of it — that name is private to that module, matching this codebase’s usual choice for a one-line filter (AGENTS.md section 5; `optimize.py`’s `_movable_disks()` is the same precedent). And `_live_transient_check()` re-derives the section 8.1 reserve formula against the target’s live `storage_status()`, not the in-memory model. Either kind of mismatch stops the group’s run with “replan\_needed” rather than trying to patch the plan around it — IMPLEMENTATION\_PLAN.md section 9.2 is explicit: “abandon the remaining moves... do not attempt to patch it.”

`_live_transient_check()`’s arithmetic is the identical section 8.1 formula `schedule.transient_invariant_ok()` checks against the in-memory model, both now calling the same `reserve.transient_charge_ok()` (see 95-schedule.md) — this one re-typed against `PveClient.storage_status()`’s live `used/total` instead of a summed disk list. The same *rule*, a different *data source*, not a second implementation of it (AGENTS.md section 5). `transient_charge_ok()` already takes a *list* of charges, not one disk, so it is ready for a future concurrent executor to call with every move currently in flight against the same target — this module does not do that yet (see “What `execute_plan()` deliberately does not do” below).

### Why “done” needs three conditions, not one

IMPLEMENTATION\_PLAN.md section 9.3.2’s completion criterion is deliberately stronger than “the task succeeded”:

```
move m from a to b is DONE ⇔ task(upid) exitstatus == OK
    ∧ volume(m) absent from GET /storage/{a}/content
    ∧ config(vmid(m)).lock is empty
```

A storage with `saferemove` enabled keeps a storage-level lock on the *source* for as long as it takes to zero the old volume at `saferemove_throughput` (10 MiB/s by default) — hours or days for a large disk, entirely after the `move_disk` task itself has already reported success. `_wait_for_move_completion()` polls the task first (unbounded — a move that was ever scheduled already passed `migration.max_single_move_duration` at planning time, so there is no separate timeout for the task loop itself), then, only when `execution.source_release.wait` is true and the source actually `saferemoves`, polls the source’s own content listing and the VM’s lock together, bounded by `execution.source_release.timeout_seconds`.

Hitting that bound is not a failure — it is "draining": the mirror is done, the wipe is still running, and the next run will see the storage as it actually is.

`execute_plan()`'s own `largest_by_storage` bookkeeping (section 8.1's `Z_s`, the largest resident disk on a storage — needed for *later* moves in the same run that check the same target's live transient invariant) updates on "moved" **and** "draining", not "moved" alone: the mirror itself is physically complete either way, so a target storage's own accounting must already include this disk regardless of whether its *source* has finished releasing.

A "draining" outcome also adds that move's *source* to a `drained_storages` set `execute_plan()` carries for the rest of the run (REVIEW.md S-05): the next move whose own source or target is in that set is skipped outright, `_drained_skip_outcome()` reporting it without a single API call. Section 9.3 asks for exactly this ("exclude it as both source and target for the remainder of the run") — a draining storage holds a storage-level lock for as long as the wipe runs, so a subsequent `move_disk` touching it would either queue behind a potentially day-long wipe (the task-status poll loop is unbounded) or fail outright, tripping `abort_on_failure` for no reason a human would consider a real failure. An earlier revision of this module tracked the (C4) *accounting* consequence of a draining target (the paragraph above) but never actually excluded the storage from later moves in the same run — found by a review pass, not a test, since nothing at the time exercised two moves sharing a storage within one run.

### VM locks are an open set, waited out, never whitelisted

`_wait_for_unlocked()` treats any non-empty `config.lock` as "wait," full stop — never a lookup table of "harmless" values to skip past (`.agents/domain-invariants.md` rule 5: a future PVE version can add a lock value this codebase has never seen, and guessing wrong risks a half-completed operation on someone else's backup or snapshot). What happens after `execution.locks.wait_timeout_seconds` depends on `execution.locks.on_timeout`:

- "skip" (the default) — report the move "skipped", continue with the rest of the plan.
- "abort" — report it "failed" and set `MoveOutcome.always_stop`, which `execute_plan()`'s stop condition checks *independently* of `execution.abort_on_failure`. This was a real design bug caught by re-reading the manual's own wording before any test existed: the manual describes `on_timeout: abort` as "abort **the run**," a stronger and unconditional promise, distinct from `abort_on_failure`'s general "stop after any plain task failure" policy, which must still be respected on its own terms for an ordinary `move_disk` failure. Folding both into one flag would have made `abort_on_failure: false` silently defeat an operator's explicit `on_timeout: abort` choice.

### The injectable clock

Every wait loop in this module measures elapsed time via `clock.now()`, never by counting `clock.sleep()` calls — so a test's fake clock, whose `sleep()` advances its own `now()` instantly instead of blocking (`.agents/testing.md`: "no real time.sleep; inject the clock"), exercises the *exact* timeout arithmetic production does, at zero wall-clock cost. `_REAL_CLOCK` is a module-level singleton (`Clock` holds only function references, never per-call state, so sharing one instance is safe) used as `execute_plan()`'s default — a bare `Clock()` there would be a flake8-bugbear B008 violation (a function call in a default expression).

### Orphaned volumes: reported, never deleted

`_detect_orphan_volumes()` runs only after a "failed" outcome, cross-referencing the target storage's content listing against every volume the VM's config actually still references (via the shared `topology.parse_disk_spec()/topology.DISK_KEY_RE` — the same disk-spec parser `topology.py` itself uses, not a second one). A volume owned by the VM but not referenced anywhere in its config is reported in that move's own `MoveOutcome.orphaned_volumes` and logged at warning — never removed automatically (`.agents/domain-invariants.md` rule 6). A `PveApiError` while checking (the storage or the VM briefly

unreachable right after a failure) is caught, logged, and degrades to an empty tuple rather than turning an already-failed move into a crash.

### confirm's callback lives in cli.py, not here

execute.py's ConfirmCallback type is Callable[[ScheduledMove], str] returning "y"/"n"/"a"/"q" — execute\_plan() calls it, interprets the four answers, and raises ValueError on anything else, but never itself talks to a terminal. cli.py is the only module allowed to call print()/input() (its own module docstring), so cli.\_confirm\_move\_interactively() is where the actual prompt text lives and where an out-of-set answer is *retried* rather than turned into the ValueError execute\_plan() would otherwise raise — a typo at the keyboard should not look like a bug report.

### What execute\_plan() deliberately does not do

- **No re-plan loop inside this module.** A "replan\_needed" outcome stops the group's run cleanly here; re-invoking the whole gate/solve/schedule/payback pipeline from the newly observed state is cli.py-level orchestration across multiple execute\_plan() calls (cli.\_run\_auto\_group(), auto mode only — see below), not something this module does itself. dry-run/confirm never re-plan at all: the operator re-runs apply by hand once ready.
- **Section 7.3's saturation check is planning-time only, not execution-time, concurrently or sequentially.** The defer check itself is enforced before either executor ever sees a flagged move — it is excluded from schedule\_result.order at planning time (docs/internals/96-payback.md) — but neither executor re-checks the ceiling *during* execution against the live in-flight set, so section 8.1 point 4 of concurrency\_ok (summing  $\omega_{\text{role}}$  over everything actually in flight right now, mirroring or draining) stays a documented gap the concurrent executor inherits rather than closes.
- **Crash recovery is not this module's own job.** execute\_plan() accepts on\_inflight\_started/on\_inflight\_finished callbacks (execute.InflightCallback) and calls them around each move\_disk — writing what they actually persist, and section 13's startup scan that reads it back, both live in crashrecovery.py/cli.py. See [93-crashrecovery.md](#).

### cli.py: one planning pipeline, shared by plan and apply

\_plan\_group() is \_handle\_plan()'s former per-group loop body, extracted so \_handle\_apply() calls the *same* function rather than re-deriving “gate, then solve, then schedule, then payback” a second time (AGENTS.md section 5) — IMPLEMENTATION\_PLAN.md section 9.2's own re-plan protocol says to re-run exactly this pipeline, not a parallel one apply keeps to itself. \_GroupPlan carries every stage's result (or None for a stage that never ran because an earlier one stopped short, e.g. load\_error set or the gate said NO ACTION) so both handlers can tell those cases apart without guessing from missing dict keys.

\_handle\_apply() acquires state.py's advisory lock **before** touching PVE or Prometheus at all — a second concurrent apply invocation (any mode, including dry-run) exits 0 immediately without paying for either round-trip first, matching section 11.2's “a live PID means another instance is running: exit 0 quietly.” Because an operator's [q]uit (or a replan\_needed, though that one is per-group and does not otherwise propagate) can stop the whole run before every group is even planned, \_render\_group\_plan\_human()/\_render\_group\_plan\_json() — the shared per-group renderers behind both plan's and apply's output — look up gate\_decisions with .get(), not []: plan always plans every group first, but apply genuinely can leave a later group completely unvisited, which must render as “not evaluated this run,” not a KeyError. This was caught by a test exercising the quit path, not by inspection.

After a group's execution, \_record\_executed\_moves() decides which moves counted as “executed” for state.json's own bookkeeping — “moved” and “draining” (the mirror itself completed either way, mirroring execute.py's own (C4) accounting choice for those same two statuses), never “would\_move”/“skipped”/“failed”/“replan\_needed”. It reads straight off the final ExecutionResult.outcomes,

never a `_GroupPlan`'s own `schedule_result.order`: a payback refusal (S-02) can already make outcomes a reordered subset of that order, and in auto mode a re-plan (`_run_auto_group()`, below) can execute moves from a *later* plan attempt that never appeared in the group's first-attempt order at all — an earlier revision iterated the stale order and silently dropped those later moves' cooldowns. For every move that counts, it calls `state.with_recorded_balance()` (this group's measured load vector, feeding the next run's drift gate) and `state.with_recorded_cooldown()` — a fresh disk cooldown for every executed disk at its *new* location, and a fresh storage cooldown for **both** of the move's storages, source and destination (section 6: "a storage involved in a migration ... accepts no new incoming moves" covers both; REVIEW.md S-03 — an earlier revision of this code recorded the destination only, which meant no configured gates. `cooldown_per_storage` could ever protect a still-draining *source*, exactly the scenario section 9.3's knob-sizing rule and `verify-storages`' own warning describe). Recording both endpoints is independent of `heuristic.py`'s own *enforcement*, which stays destination-only: a cooldown storage is excluded as a move/swap target, never as a source a disk may still leave (`docs/internals/90-heuristic.md`). The accumulated state is written once, at the very end of the whole run (every group, not per-group), through `state.save_locked_state()` — never `state.save_state_atomic()`'s `rename`, which would silently detach the very lock `acquire_lock()` just took (see `docs/internals/15-state.md`'s account of that bug) — and only then released via `state.release_lock()`, in a finally so a mid-run exception never leaves the lock held.

### The payback gate: apply refuses on its own verdict (REVIEW.md S-02)

Section 7 calls the payback rule "a **hard acceptance test on the finished plan**, not merely a soft  $\gamma$  penalty" — a verdict plan shows the operator, not a suggestion. An earlier revision of `_handle_apply()` computed `payback_result` and then ignored it entirely on the execution path: a plan whose aggregate economics failed, or whose one move exceeded `migration.max_single_move_duration` (`payback_result.rejected_moves`, the *hard* per-move rule §7.3 separates from the aggregate one), was executed anyway. `_apply_payback_gate()` is the fix: it never calls `execute_plan()` at all for a move in `rejected_moves`, or for *any* move in the group when not `payback_result.aggregate_ok` — both report a "skipped" outcome ("refused: ...") through the exact same `MoveOutcome/ExecutionResult` shapes `execute.py` itself produces, so neither the renderers nor `state.json`'s bookkeeping need a special case: a refused move was never "executed" (it is not "moved"/"draining"), so it earns no cooldown and contributes nothing to `last_balance`. A plan that clears both checks is passed to `execute_plan()` with its individually-rejected moves (if any) simply absent from `order` — which is also what makes `_wait_for_move_completion()`'s own docstring claim ("a move that was accepted at planning time already passed `migration.max_single_move_duration`") true by construction rather than merely asserted. In confirm mode, the group's payback lines print once, before the first prompt — an operator confirming moves one at a time sees the tool's own verdict before being asked anything, not only in the post-run report.

### auto mode: timewindow.py and \_run\_auto\_group() (phase 8)

`execute_plan()` itself grew two auto-only budgets, both `None` (inactive) unless `cli.py` passes real values: a deadline (an absolute instant, checked against `clock.now()` before starting each move, together with that move's `payback.MoveCost` duration estimate from `move_costs_by_key`) and a `max_migrations` count, decremented once per move actually attempted ("moved"/"draining"/"failed" — never "skipped"/"would\_move"/"replan\_needed"). Both stop the whole run cleanly (never mid-move) via the same `_auto_budget_stop_outcome()` helper, since remaining time only ever decreases and a spent cap stays spent — there is never a reason to skip one move for a budget reason and then try a later one anyway.

The deadline is actually checked **twice** for a locked VM: `_auto_budget_stop_outcome()`'s pre-flight check runs before any network call for this move, but a VM-lock wait (`_wait_for_unlocked()`) can burn a large, unpredictable share of the same budget — up to `execution.locks.wait_timeout_seconds`, hours

by default — entirely *after* that first check already said yes. `_execute_one_move()` re-checks the identical condition (`_deadline_exceeded()`, the one shared implementation of “does this much more time still fit”) right after a lock clears and before committing to `move_disk`, refusing the move there instead if the wait alone consumed what was left. Caught by re-reading section 9.1 against the code before any test existed, not by a failing test.

`timewindow.py` is a small, dependency-free module answering exactly one question, `current_deadline(windows, now) -> datetime | None`: `None` means no restriction at all (no `time_windows` configured), and a non-`None` value is either the close time of whichever configured window covers `now`, or `now` itself when none currently do (zero remaining budget — deliberately not a separate “not in any window” signal, so a caller’s “does the next move fit before the deadline” arithmetic refuses correctly without a second check). It works in **local time**, not UTC — a documented, deliberate choice (the plan does not name a timezone for this setting) matching `systemd.timer`’s own `OnCalendar=` default and how an operator actually thinks about a maintenance window. The classic overnight-window bug — a days: [fri] window tagged Friday still needing to match at 2am *Saturday* — gets its own two-piece matching logic (`is_window_active()`) and its own regression tests.

`timewindow.py` itself never reads the clock (every function takes `now` as a parameter); `cli._real_local_now()` is what production actually passes. It resolves the host’s real IANA zone from `/etc/localtime`’s own symlink target (the standard mechanism on Debian and virtually every Linux distribution) rather than settling for `datetime.now().astimezone()`’s *fixed*-offset result — the difference matters specifically for `window_close()`’s “closes tomorrow” case, which combines today’s `tzinfo` with tomorrow’s date: a real `zoneinfo.ZoneInfo` re-resolves its own UTC offset for that date, a frozen offset does not, so an overnight window could otherwise stay “active” up to an hour longer than configured on the two nights a year the local zone’s clocks actually change. `_real_local_now()` falls back to the frozen-offset form only when the zone name cannot be determined at all, at which point that same narrow, two-nights-a-year discrepancy returns — a residual, documented limitation, not a silent one.

`cli._run_auto_group()` is where section 9.2’s re-plan protocol actually lives: a loop, bounded by `execution.max_replans_per_run`, that calls `_apply_payback_gate()` and then, only if the *last* outcome it produced was “replan\_needed”, re-fetches topology fresh (`build_topology()` — not just re-running `_plan_group()` on the same `Group` object, since whatever triggered the mismatch can mean the group’s own membership changed) and calls `_plan_group()` again before retrying. A re-plan that concludes `NO ACTION` (or hits a load error) ends the loop quietly, not as a failure — section 9.2’s own words, “the gates may well conclude no further action is needed.” Exceeding the replan cap rewrites the final `stop_reason` to name the setting explicitly, rather than reporting the *symptom* (the last mismatch) as if it were the cause.

One deliberate simplification: everything the human/JSON report shows for a group (`gate_decisions[name]`, `schedule_results[name]`, etc.) still comes from the group’s *first* planning attempt, even after auto re-plans it — only the accumulated `ExecutionResult.outcomes` (every attempt’s outcomes, concatenated) and `last_balance`’s recorded load vector reflect what actually happened. Returning a second `_GroupPlan` from a successful re-plan would need `_handle_apply()`’s reporting dicts to cope with the *first* attempt’s fields being live while a later one supersedes them, and — worse — a re-plan that lands on `NO ACTION` would have to report a group whose `schedule_result/payback_result` are `None` despite the group having `ACTed` and moved something moments earlier. Reporting the first attempt’s plan alongside the full outcome history is simpler and does not lose information: the report already shows *why* it stopped and *what happened next* per outcome.

`execution.max_migrations_per_run` is a per-*invocation* budget, shared across every group `_handle_apply()` visits — `_run_auto_group()` returns the updated remaining count, which its caller threads into the next group’s call rather than resetting it.

## Concurrent execution: `_execute_concurrent()` and its non-blocking helpers

`execute_plan()` dispatches to `_execute_concurrent()`, an entirely separate orchestration loop from `_execute_sequential()` above, whenever `mode == "auto"` and either `execution.max_concurrent_migrations` or `max_concurrent_per_storage` is configured above its default of 1. Every other case (dry-run, confirm, or auto at the default caps) uses `_execute_sequential()` completely unchanged — this is a genuine regression guarantee, not just a claim: every test written before concurrency existed still passes unmodified, because none of them touch the new dispatch condition.

**Why a second loop, not one generalized loop.** The sequential loop's `_wait_for_unlocked()` and `_wait_for_move_completion()` both *block*, sleeping in a tight loop until their own condition resolves. A concurrent executor cannot reuse either directly: blocking on one move's lock or completion would stall every *other* in-flight or candidate move for as long as that wait lasts, which defeats the entire point of running several at once. Both were refactored into a non-blocking step plus a thin blocking wrapper *before* concurrency was added at all, so the refactor's own correctness could be verified by the full pre-existing test suite passing unchanged, independent of any new concurrent-specific test:

- `_check_lock_once()` is the one live lock read; `_wait_for_unlocked()` is now just that read in a loop.
- `_poll_move_once()` is section 9.3.2's completion criterion as a single step, returning `None` while still in progress or `(status, detail)` once terminal; `_wait_for_move_completion()` is now just that in a loop. `_MoveWaitState` carries the one thing that has to survive across calls (`drain_start`, once the source-release phase begins) since there is no longer a local variable's lifetime to hold it.

`_launch_decision()` is section 9.2's five pre-flight re-checks plus section 8.1's generalized transient invariant, for one poll cycle, without blocking — the concurrent counterpart to `_execute_one_move()`'s own pre-flight-then-lock-then-live-check sequence. `_LockWaitTracker` plays the same role for a lock wait that `_MoveWaitState` plays for a move's completion: start/warned survive across cycles so the executor logs the "VM is locked" warning once, not once per poll, and knows when `execution.locks.wait_timeout_seconds` has actually elapsed. The generalized invariant itself lives in `reserve.transient_charge_ok()` (`docs/internals/95-schedule.md`) — `_launch_decision()`'s only job is to assemble its inputs: a live `storage_status()` read for used/total, this run's own (C4) `largest_by_storage` tracking for `Z_b`, and `_target_charges()` (every other in-flight move's own disk size, for moves already landing on the same target) alongside the candidate's own size.

**Strict FIFO, deliberately.** `_advance_pending()` only ever considers `pending[0]` — the same move a sequential run would try next. If it cannot launch yet (a per-storage cap, or a lock that has not cleared), the executor does not skip ahead to a later candidate that might launch immediately; it waits (polling everything already in flight in the meantime) and re-evaluates the same head next cycle. Several moves genuinely run concurrently once launched — nothing ever blocks waiting for one to finish before starting another — but *which* move launches next is always the scheduler's own next-in-line choice. A more sophisticated scheduler could reorder around a blocked head to keep every slot busy; this one can only ever under-deliver on throughput doing that, never launch a move out of the payback-scored order `schedule.order_moves()` computed, matching this codebase's existing tolerance for a documented, safe simplification over a more optimal but harder-to-verify one (`schedule.py`'s own deferred ordering-priority-2 and staging are the same trade).

**A stop condition drains, it does not abandon.** Once a budget is exhausted, a move fails under `execution.abort_on_failure`, or a pre-flight/live-invariant mismatch produces "replan\_needed", `_execute_concurrent()` stops calling `_advance_pending()` (no further moves launch) but keeps polling whatever is still in flight to its own natural conclusion before returning — their outcomes, and their crash-recovery callbacks, must still fire. Whatever is still in pending at that point (never even pre-flighted) is abandoned unreported, exactly like `_execute_sequential()`'s own early return already leaves every

move after the one that triggered a stop completely unaccounted for.

### Crash and two-instance recovery: crashrecovery.py

**What does this page answer?** How does an apply run notice a move\_disk left running by a crashed previous run, or by a second instance running concurrently; how does state.json's inflight\_upids actually reach disk *before* the move it names can crash the engine; and what does a discovered vmid do to the rest of the run? Describes proxmox\_storage\_drs/crashrecovery.py, state.py's with\_inflight\_upid()/without\_inflight\_upid(), and the execute.py/cli.py wiring around them.

### The two failure modes section 13 names, one mechanism

"Engine crashes mid-move" and "Two DRS instances running" are different causes with the same shape: a move\_disk task keeps running against a VM this run does not know about. state.json's inflight\_upids field exists so a run notices either case *before planning anything* (section 13's own words) rather than proposing a second, conflicting move against the same disk:

1. **A local trace.** execute.py writes a UPID into inflight\_upids immediately after move\_disk() returns for it, and clears it once that task itself has finished — see "Writing before the crash" below. A process that dies mid-move (killed, OOM, host reboot) leaves that UPID behind; the *next* run's startup scan finds it.
2. **A cluster-wide scan.** state.json's own fcntl.flock() (docs/internals/15-state.md) is real, kernel-enforced mutual exclusion for two instances on the *same host*, but cannot see another node at all. PveClient.cluster\_tasks() (GET /cluster/tasks) is the one read that actually crosses the cluster — crashrecovery.py scans it for a still-running qmmove task from this tool's own configured user that the local trace does not already know about. IMPLEMENTATION\_PLAN.md is explicit that this is what degrades "two instances running" from "conflicting" to merely "slow and redundant" — the correct fix is still "run the timer on exactly one host" (cron/systemd configuration, not this code); this is the safety net for when that is not honoured.

Either way, reconcile\_inflight() reports the affected vmid for cli.py to exclude from this run's planning — **never** force-cancelled, never assumed safe to ignore.

### parse\_upid(): PVE's UPID grammar, confirmed live

A UPID is UPID:{node}:{pid}:{pstart}:{starttime}:{type}:{id}:{user}:. id is the vmid, as a string, for a VM-scoped task type (qmmove among them); it is empty for a cluster-scoped one (vzdump, say). This shape is not general PVE-API knowledge taken on faith — .agents/domain-invariants.md rule 10 forbids stating unverified PVE behaviour — it was confirmed live against the real dev cluster ([dev-cluster-access]]) by reading an actual qmconfig task's own UPID before parse\_upid() was written to depend on it. parse\_upid() returns None for anything that does not parse as a UPID at all (a foreign or malformed string), never raises — a scan finding one unparseable entry must degrade to "cannot use this one," not abort.

### /cluster/tasks' own convention

PveClient.cluster\_tasks() (GET /cluster/tasks) returns every node's recent/active tasks as {node, type, id, upid, user, starttime, saved}, plus endtime/status — **both present once the task has finished, both absent while it is still running**. This is a different convention from the per-task PveClient.task\_status() endpoint (GET /nodes/{node}/tasks/{upid}/status), whose own status field is the string "running" or "stopped" — confirmed live against the same cluster alongside the UPID grammar above, and the reason\_scan\_cluster\_for\_foreign\_inflight() uses "endtime" in task or "status" in task to filter obviously-finished candidates cheaply, then \_task\_is\_running() (below,

wrapping `task_status()` to actually confirm one that looks live — never the reverse, and never treating `/cluster/tasks`' own status string as if it used `task_status()`'s "running"/"stopped" vocabulary.

### **reconcile\_inflight(): both halves**

```
def reconcile_inflight(
    client: PveClient, state: State, auth: AuthConfig
) -> tuple[frozenset[int], State]:
```

Split into two helpers purely to stay within this project's flake8 complexity limit:

- **\_reconcile\_recorded\_upids()** re-checks every UPID state. `json` already remembers via `_task_is_running()` (which wraps `task_status()`, not `/cluster/tasks`' own convention — see above). One whose task has since finished is dropped from the returned `inflight_upids` (an ordinary reconciliation of a move the engine never got to clear itself, logged at `INFO`); one still running is kept and its `vmid` excluded. A malformed or `vmid`-less recorded entry is dropped outright — there is no future run in which it could ever resolve either.
- **\_scan\_cluster\_for\_foreign\_inflight()** is the cluster-wide half described above, skipped entirely (`frozenset()`, no API call at all) when `expected_task_user()` cannot determine this tool's own configured user (neither `token_id` nor `username` set — config validation makes this unreachable in practice, but the function degrades safely rather than guessing). `already_known` (the still-inflight subset of `state.json`'s own recorded UPIDs) is subtracted first, so a run's own in-flight move never logs a spurious "found a foreign task" warning about itself.

`_task_is_running()` returns `None`, not an exception, when the check itself fails (the node is unreachable, say) — logged at `WARNING` and treated by both callers as "assume still running." That is the safe direction to guess wrong in: excluding a VM that turned out to be free costs one skipped move on this run; assuming free a VM that was not risks a second move `_disk` racing the first.

### **Writing before the crash: `execute.py`'s callbacks**

Persisting `inflight_upids` only at the *end* of a whole apply run — the way `last_balance/cool downs` are persisted — would never survive the exact crash this mechanism exists to recover from: if the engine is killed mid-move, nothing after that point ever runs. `execute_plan()` therefore accepts two optional callbacks, `execute.InflightCallback = Callable[[str], None]`:

- **on\_inflight\_started** fires in `_execute_one_move()` immediately after `move_disk()` returns a UPID — before anything else happens with it, including the (potentially very long) wait for it to finish.
- **on\_inflight\_finished** fires immediately after `_wait_for_move_completion()` returns, for *any* outcome including "failed" and "draining". A "draining" source is no longer tracked by this UPID at all once its task itself has finished — only by its own content-listing poll (`state.without_inflight_upid()`'s own docstring) — so it is cleared here exactly as promptly as a "moved" outcome, not left recorded until the source actually releases.

Neither call is wrapped in `try/finally`: if something inside the wait itself raises (a network failure mid-poll), `on_inflight_finished` simply never fires and the UPID stays recorded — which is correct, since the move might still be running and the next startup's scan must still find it.

`cli.py`'s `_make_inflight_callbacks()` is what the two callbacks actually do: write straight to disk via `state.with_inflight_upid()/without_inflight_upid()` followed by `state.save_locked_state()`, through the same `LockHandle._handle_apply()` itself holds for the whole run. This is the one deliberate, narrow exception to this codebase's otherwise-pure functional state-threading style (`_handle_apply()`'s state variable is reassigned throughout, never mutated) — `_InflightStateBox` is a mutable box purely so a closure passed several call-frames deep into `execute_plan()` can still reach it. `_handle_apply()`'s

own finally block saves `state_box.value`, never the plain state variable, for exactly this reason: if an exception propagates out of a group's execution, a callback earlier in that same group can already have written a newer `inflight_upids` straight to disk than whatever state was last assigned in the surrounding loop, and saving the stale local variable would silently clobber that write.

### **cli.py: the startup scan, before planning anything**

`_handle_apply()` calls `_reconcile_inflight_and_fold_exclusions()` immediately after building the PVE client, before `build_topology()` for planning — section 13's own “before planning anything.” A non-empty result is folded into `exclude.vmdids` on the `ResolvedConfig` used for the rest of this run, reusing `topology.py`'s existing (C2) `vmid/tag` pin rather than a second exclusion mechanism (AGENTS.md section 5: one implementation); the report shows such a `vmid` pinned with the generic “excluded by config (`exclude.vmdids`)” reason, and `crashrecovery.py`'s own `logger.warning()` calls are what tell the operator why it is there. This runs for every mode, including dry-run — an accurate dry-run report should already reflect a `vmid` a real apply run would also have excluded, and reconciling a since-finished recorded UPID away is harmless (and useful) even when nothing gets executed this run.

### **Ordering the moves**

**What does this page answer?** How does `schedule.py` turn a target assignment into an ordered, transient-feasible sequence of moves, and what happens when no order exists? Describes `proxmox_storage_drs/schedule.py`.

### **One loop, reusing heuristic.py's objective for the ratio**

`order_moves()` implements section 8.2's pseudocode directly: while there are pending moves, find the ones that are transient-feasible right now, schedule the best by imbalance-reduction-per-byte (moves that resolve a currently-violating storage win outright, regardless of that ratio), apply it, repeat. “Imbalance reduction” is computed by calling `heuristic.evaluate_assignment()` twice (before/after the candidate move) and diffing `.imbalance_term` — not a second formula, the same one `heuristic.py` already computes, so a plan's ordering agrees with whatever `objective.spread_metric` is configured exactly the way the assignment that produced it did.

### **cost\_m is z\_d — and that is not an approximation today**

Section 8.2's ratio is “imbalance reduction / `cost_m`”. `cost_m` here is simply the disk's size in bytes, not `payback.py`'s real duration-based cost (mirror time plus, when `saferemove` is on, wipe time — and `wipe throughput` is a per-storage config value, so two candidate moves off different sources can have genuinely different wipe costs even for the same `z_d`). `migration.bwlimit_bytes_per_sec` is one global config value, so every move's *mirror* duration alone is `z_d / bwlimit`, a constant divided into every disk's size equally — dividing an imbalance reduction by `z_d` and dividing it by `z_d / bwlimit` produce the **same ordering** (`bwlimit` is a positive constant common to every candidate). So `cost_m = z_d` reproduces the true mirror-only ordering exactly, but is an approximation once a group has moves whose `saferemove` wipe cost differs enough to change the ranking `payback.py`'s full cost would produce. Scoring candidates by `payback.compute_move_cost()` instead of `z_d` is a natural follow-up, not yet done: this module predates `payback.py` and has not been revisited to consume it (REVIEW.md R-04).

### **The transient invariant, called with a single-move set always**

Section 8.1 defines a generalized invariant over a whole in-flight set  $M$  specifically so it can be “implement[ed] as the single feasibility predicate and call[ed] with  $M = \{m\}$  for the sequential case” — this module is that sequential case, always. `transient_invariant_ok()` checks one move against state, the assignment as of immediately before that move starts; every previously-scheduled move is assumed **fully complete** by then (mirror finished, and any `saferemove` wipe finished, source genuinely freed) before

this one begins. That models `schedule.order_moves()`'s own job correctly regardless of `execution.max_concurrent_migrations` — it always computes one, strictly *-sequential order*; that order becomes a launch queue several moves can run against concurrently only once `execute.py`'s own executor gets hold of it (92-execute.md's "Concurrent execution" section) — this module itself never reasons about overlapping in-flight windows, and does not need to, since `execute.py`'s own live re-checks (below) are what actually enforce the invariant against whatever really ends up running together.

The actual arithmetic —  $\text{used}_b + \text{sum}(z_m) + f_b * \max(Z_b, \max(z_m)) \leq C_b$  — lives in `reserve.transient_charge_ok()`, taking a *list* of charges rather than one disk, so it degenerates to section 8.1's original single-move form when called with `[disk.size_bytes]` (what this module does) and generalizes correctly to several moves landing on the same storage at once when `execute._launch_decision()` calls it with more than one, for real, under concurrent execution (AGENTS.md section 5: this is the *one* place that formula is written, not two functions that happen to agree). `execute._live_transient_check()` (sequential) and `execute._launch_decision()` (concurrent) both call it against a live `storage_status()` read instead of this module's model-derived numbers — see 92-execute.md.

The `min_free_bytes` floor is folded into the transient check the same way (C5) folds it into the steady-state one ( $\max(f_b * \max(Z_b, z_d), \text{min\_free\_bytes})$ ) — the plan's own section 8.1 formula does not mention the floor, but there is no reason a storage's absolute minimum free space should stop applying just because a migration happens to be in flight.

### **The heuristic can accept a residual violation; the scheduler cannot execute one**

`heuristic.py`'s repair step can leave a storage still violating (C5) if no reachable assignment can fully clear it — a genuine, sometimes unavoidable outcome the objective's `reserve_penalty_term` reports rather than hides (see 90-heuristic.md). When `order_moves()` is asked to schedule the move that leads to exactly that residual state, the transient check correctly says no: landing a disk on a target that still would not satisfy (C5) is not transient-feasible, full stop, and this module never treats "the heuristic already decided this was the best available" as a reason to relax that. The result is a deadlock report for that move (section 8.3 option 3, "report... never force a move that breaches the reserve") — an honest "this cannot be safely executed", not a false all-clear. See `test_plan_reports_a_deadlock_when_even_the_best_target_still_violates` in `tests/unit/test_cli.py` for the exact scenario reproduced end to end.

### **final\_assignment is the schedule's own truth, not the heuristic's target**

`ScheduleResult.final_assignment` is the assignment produced by actually applying order in sequence, starting from the current placement — every disk, not only the moved ones. When nothing deadlocks it is identical to the heuristic's target assignment, but when deadlocked is non-empty (fully or, more subtly, only partially) it is not: a caller reporting an "after" utilization/spread must evaluate `final_assignment`, not `heuristic.HeuristicResult.assignment`, or it reports numbers for moves that were never actually scheduled (REVIEW.md R-02 — `cli.py`'s plan originally did exactly this, correct only when `order_moves()` scheduled every move a plan proposed).

### **What this pass deliberately does not do**

- **Reasoning about concurrency at scheduling time.** This module always produces one, strictly-*sequential order* — `execute.py`'s own executor is what actually runs several of that order's moves at once under `execution.max_concurrent_migrations > 1`, re-validating section 8.1's generalized invariant live against whatever really ends up in flight together (92-execute.md's "Concurrent execution" section). `payback.py`'s move-duration estimates would let this module choose a smarter, concurrency-aware *order* (e.g. front-loading moves that can provably run together) instead of leaving that entirely to `execute.py`'s own strict-FIFO launch discipline, but that scheduling-time optimization is not implemented.

- **Ordering priority 2** (“moves that free space a later move needs”) and **staging** (section 8.3 option 1). Both are refinements over a plan that is already feasible and safe without them — priority 2 only affects how well a plan front-loads its value if interrupted, and staging only helps in a genuine pebble-motion deadlock. Omitting them can only make this module report a deadlock in a case a more sophisticated scheduler would have found a way through; it can never produce an order this module thinks is safe but is not.
- **Cooldowns**, for the same reason `gates.py` doesn’t implement them yet: `state.json` (section 11.2), which would record the timestamps a real cooldown check needs, does not exist.
- **The section 7.3 saturation check** — belongs with `payback.py` (implemented; see 96-`payback.md` for why the check itself is deferred).

### Wired into plan, not yet into apply

`cli.py`’s `plan` command calls `heuristic.run_heuristic()` then `schedule.order_moves()` for every group the gate says to act on, and renders the result (see `docs/manual/27-plan.md`). Nothing calls `schedule.py` from `apply` yet, because `apply` itself does not exist (`execute.py`, phase 7+) — `plan` only ever computes and prints.

### Migration cost and the payback test

**What does this page answer?** How does `payback.py` turn a scheduled move into a cost, a plan into a benefit, and the two into an accept/reject verdict — and why does a reserve-fixing plan always pass? Describes `proxmox_storage_drs/payback.py`.

### Cost and benefit are computed from data other modules already have

`compute_move_cost()` needs only a `schedule.ScheduledMove` and its source topology. `Storage` — no new fetch, no new state. `duration_mirror` is  $z\_d / \text{migration.bwlimit\_bytes\_per\_sec}$ ; `duration_wipe` is  $z\_d / \text{saferemove\_throughput}$  when `migration.account_saferemove_wipe` and the source has `saferemove` on, else zero. `compute_benefit_load_seconds()` takes two `heuristic.raw_spread()` values — the pre-plan and post-plan `raw`, unweighted section 7.2 spread quantity — and multiplies their difference by `migration.payback_horizon_seconds`. **Not** `ObjectiveBreakdown.imbalance_term`: that is the same quantity scaled by `objective.alpha_spread` for the section 5.4 *solver* objective, and passing it here would make the payback ratio depend on that tuning knob rather than only on the imbalance reduction and migration cost a plan produces (REVIEW.md R-01 — earlier code passed `imbalance_term` directly; only invisible while `alpha_spread`’s default of 1.0 made the two numerically identical). `cli.py`’s `plan` handler computes `raw_spread()` for the pre-plan and post-plan assignments and passes those through.

### headroom\_src/headroom\_dst: a plan formula this project cannot fill in

Section 7.1 writes `duration_mirror_d = z_d / min(bwlimit, headroom_src, headroom_dst)`. Neither `headroom_src` nor `headroom_dst` is defined anywhere else in `IMPLEMENTATION_PLAN.md`, and no `config` field or topology. `Storage` attribute represents a per-storage effective throughput ceiling distinct from the one global `migration.bwlimit_bytes_per_sec`. `compute_move_cost()` uses  $z\_d / \text{bwlimit}$  only — exactly what the formula reduces to whenever neither storage’s own throughput is the binding constraint, which is the case `bwlimit` exists to enforce in the first place. This is recorded as a known simplification of the plan’s own underspecified formula, not silently worked around.

### The reserve-override exemption

Section 7’s payback test weighs a move’s cost against the *balance* benefit it buys. That framing has an edge it does not name: a move resolving an active (C4)/(C5) violation is not optional the way a balance-driven move is, and its “benefit” under the section 7.2 formula can easily be zero or even structurally zero — moving the only loaded disk in a two-storage group between the two storages changes which one

carries it without changing  $E$  at all, no matter how urgently the move is needed for capacity reasons. Rejecting that move on economic grounds would contradict section 13’s own “the reserve is never traded against balance,” which `gates.py` already treats as absolute for the drift/ imbalance gates. `evaluate_plan_payback()` applies the identical rule here: **a plan containing any move with `resolves_reserve_violation=True` always passes the aggregate ratio test**, regardless of the computed ratio. The hard per-move `max_single_move_duration` rule is not exempted this way — section 7.3 lists it as applying “regardless of the aggregate test” precisely because it is an operational limit (a mirror that takes that long has other costs an economic ratio does not capture), not an economic one.

`test_plan_json_output` and `test_plan_json_output_accepts_payback_when_saferemove_is_off` in `tests/unit/test_cli.py` exercise both halves of this end to end: the same reserve-driven, zero-benefit move is accepted when its duration is within the limit and rejected (correctly, via the hard rule, not the economic one) when a slow `saferemove` wipe pushes it over `max_single_move_duration`.

### The section 7.3 saturation guard: `compute_move_cost()`’s optional target

`compute_move_cost()` stays pure (the module docstring’s own promise: “nothing fetches anything”) by taking the guard’s inputs already computed, rather than fetching a forecast itself: `target`, and `l_hat_src/l_hat_dst` — the caller’s own already-computed `L_hat_s(duration_mirror)` for each endpoint (`forecast.storage_upper_bound()`, section 10.1, summed over the disks *currently* resident on that storage — not the moving disk’s own hypothetical arrival, since during mirroring it is still served from `src`, and its mirror-write traffic to `dst` is exactly what the  $\omega_{dst}$  charge below already accounts for separately). Left at their defaults (`target=None`, both  $0.0$ ) the check is simply inactive — every call site written before this existed, and `cli.py`’s own dry-run/ `plan` paths that have not been updated to compute a forecast, keep working unchanged.

`mirror_duration_seconds()` is `compute_move_cost()`’s own `duration_mirror_d` arithmetic, factored out so a caller can learn it *before* calling `compute_move_cost()` — a genuine ordering dependency: the guard’s forecast horizon is this move’s own mirror duration, but `compute_move_cost()` is also what turns that forecast into the `saturation_deferred` verdict. A caller wanting the check active must: call `mirror_duration_seconds()`, use it as the horizon for `forecast.storage_upper_bound()` against each endpoint’s own resident disks, then call `compute_move_cost()` with the results.

`_saturation_deferred()` charges `migration.source_load_weight/ target_load_weight ( $\omega_{src}/\omega_{dst}$ )` on top of each endpoint’s own `l_hat`, matching section 7.3’s  $\omega_{role}$  table for the mirroring state — the same two config values `compute_move_cost()`’s own cost formula already uses (AGENTS.md section 5: no second pair of weights invented for this). `MoveCost.saturation_deferred/PaybackResult.deferred_moves mirror exceeds_max_duration/rejected_moves`’s existing shape exactly, but are kept as distinct fields — section 7.3 itself draws the same distinction (“reject the move” vs. “defer the move to a later run”), and a report should be able to say which of the two happened, one hard and always active, the other best-effort and silently inactive wherever `saturation_load` is unset. `PaybackResult.accepted` now requires neither being non-empty.

### What this pass deliberately does not do

- **The 3-retry re-solve-with-doubled-beta/gamma loop** (section 7.3) on aggregate payback failure. A real UX refinement — it converges on the smaller subset of high-value moves rather than abandoning the run — but not a correctness requirement: reporting accept/reject plainly already satisfies “reject a plan whose cost outweighs its benefit.” Implementing the loop means re-invoking `heuristic.py/schedule.py` with adjusted weights from inside `cli.py`’s own orchestration, which is where it belongs when it is built, not half-done inside this module now.
- **Automatically dropping an individually-rejected move** and re-solving without it. Reported instead (`MoveCost.exceeds_max_duration`, `PaybackResult.rejected_moves`) — the same “report,

never force” choice `schedule.py`’s deadlock reporting already makes for an unschedulable move.

- **The section 7.3 saturation-ceiling defer check beyond its mirroring -phase reading.** `compute_move_cost()` now implements  $L_{\text{during}}(s) \leq \text{saturation\_ceiling} * N_s$  for each endpoint (see the section above), but only at the mirroring-phase horizon the section’s own header names (“push either endpoint above ... during *the mirror*”) — not a second, separate check for the *draining* phase ( $\omega_{\text{wipe}}$  over `duration_wipe_seconds`), which the full generalized in-flight-set model implies but which needs `schedule.py` to reason about overlapping moves, something it does not do (`95-schedule.md`). `N_s` unset on a storage still skips the check for that endpoint entirely, exactly as the plan’s own words allow (“fully supported... loses only this one advisory check”).

### Wired into plan and apply alike, via `_plan_group()`

`cli.py`’s `_plan_group()` (`docs/internals/92-execute.md`’s “one planning pipeline, shared by plan and apply”) computes a `PaybackResult` for every group the gate acts on, right after scheduling its moves, and both plan’s render functions and apply’s payback gate (`_apply_payback_gate()`) consume it — a rejected or deferred move never reaches `execute.py` regardless of mode. `_saturation_forecast_inputs()` is what builds the saturation guard’s own forecaster and per-disk load history, once per group, only when at least one of its storages configures `saturation_load` — skipped entirely otherwise, so a cluster that never sets it pays no Prometheus cost for a check it cannot use. `_compute_one_move_cost()` then derives each move’s own `l_hat_src/l_hat_dst` (`forecast.storage_upper_bound()` over each endpoint’s *currently* resident disks, at that move’s own mirror duration) before calling `compute_move_cost()`. `compute_wipe_duration_seconds()` is also used by `verify-storages`’s existing wipe-time warning — one implementation of the formula, not two (`AGENTS.md` section 5).