

Proxmox Storage DRS -- Internals

How the tool works, for whoever changes it

Bernd Zeimet bernd@bzed.de

2026-09-05

Rendered from docs/internals/*.md, SHA-256 41b38a2601ff3a3f060f0e0d2bd36654f844e930dc0f09517d4a58b128f54148
(sha256sum --check docs/internals.pdf.sha256 in the repository must succeed).

Copyright © 2026 Bernd Zeimet. Licensed under the GNU Affero General Public License, version 3 or later.

Contents


The pipeline, as specified and as built	4
Module layout, as it stands	4
Why config.py depends on forecast.py	5
Where to read next	5
How a YAML file becomes a validated Config	6
Resolution order	6
The two-stage validation	6
Secrets from the environment	6
What is deliberately <i>not</i> checked here	7
ResolvedConfig: what the caller actually gets	7
Forecasting: the protocol and its three implementations	7
The one rule, in one place	7
Why the upper bound, never the point estimate	7
QuantileForecaster	7
SeasonalNaiveForecaster	8
HoltWintersForecaster	8
build_forecaster()	8
The Prometheus client and verify-metrics	8
_SessionLike / _ResponseLike: why not requests.Session directly	8
_get(): one request path, three endpoint shapes	8
PromQL construction	9
verify_metrics(): six checks, six functions	9
What cli.py does with the report	9
Command dispatch, the mode-override rule, and why logs go to stderr	9
One parser, global options first	9
Command handlers: a dict, not a chain of ifs	10
The --mode escalation rule	10
Why logs go to stderr, not stdout	10
JsonFormatter: what ends up in one log line	10
--manual: preferring man(1), falling back only when necessary	11
The Proxmox VE API client	11
Why proxmoxer	11
build_client(): the one place auth is assembled	11
PveClient: one method per section 3.5 endpoint	11
move_disk(): the one and only bytes/s -> KiB/s conversion	12
Building the disk/storage/group model	13
D means every disk this module places, pinned or not	13
One pass, in the order section 3.5 lists	13
The lock field comes from vm_config(), not a separate call	14
Pin priority: _pin_reason()	14
Sizes: content is authoritative, config is the fallback	14
U _s e _x t: everything not referenced	14
reserve.py: one (C4)/(C5) evaluator, shared	14

The load model	15
One function, one group, already-built input	15
Six queries, not six-times-groups queries	15
Coverage rejection excludes a disk from the group total, not just from its own l_d	15
<code>tpmstate0/unusedN</code> are never rejected; <code>efidisk0</code> is	15
$T_g = 0$ is “idle”, computed only from accepted disks	16
L_s/u_s are the <i>current</i> assignment, not a candidate one	16
What is still missing from phase 3	16
Gating: deciding whether to act at all	16
One function, three gates, in the order section 6 lists them	16
<code>last_load=None</code> everywhere in this codebase, for now	17
<code>show-load</code> ’s gate line is diagnostic, not a real decision	17
Cooldowns are not this module’s job	17
The heuristic solver	17
<code>evaluate_assignment()</code> is the objective, kept separate from the search	17
Repair is unconditional, not weight-driven	17
Descend explores swaps, not only single moves	18
<code>objective.spread_metric</code> : two different quantities, not one rescaled	18
Proof this reproduces the plan, not just itself	18
What this pass deliberately does not do	18

What does this page answer? Where does an invocation of pve-storage-drs go, which module owns which piece, and how much of IMPLEMENTATION_PLAN.md's section 2 architecture actually exists right now?

The pipeline, as specified and as built

IMPLEMENTATION_PLAN.md section 2 describes seven stages: collect, join, gate, solve, cost, order, execute. As of this page, stages 1 (collect), 2 (join), 3 (gate) and the heuristic half of 4 (solve) exist; the MILP path, cost (payback) and order (scheduling) through execute are IMPLEMENTATION_PLAN.md section 12 phases 4 (remainder)-9 and are not yet written. Do not take this page as a claim that the whole pipeline runs end to end — [../manual/30-safety-and-status.md](#) is the authoritative per-command status: gates.py's verdict is real and show-load prints it, but nothing yet acts on it (no plan/apply), and every call site passes last_load=None since state.json (the real drift history) does not exist yet — see [80-gates.md](#). heuristic.py can already produce a target assignment for a group, but nothing yet orders its moves under section 8's transient reserve invariant or calls it from cli.py — see [90-heuristic.md](#).



Module layout, as it stands

Module	Responsibility	Plan section
exceptions.py	The DrsError hierarchy every deliberate failure raises	AGENTS.md section 5
units.py	Duration/size parsing ("24h", "200MiB") and human-readable formatting	section 11

Module	Responsibility	Plan section
config.py	Load /etc/pve/drs.yaml, jsonschema + semantic validation, the frozen dataclass config model	section 11
config_schema.json	The jsonschema structural half of validation	section 11.1
forecast.py	The Forecaster protocol, required_range_seconds, and quantile/seasonal_naive/holt_winters	section 10
logging_setup.py	Structured JSON logging to stderr	section 2.1 (amended, see 40-cli-and-logging.md)
metrics.py	PrometheusClient, PromQL construction, verify_metrics(), compute_disk_coverage()	sections 3.1-3.4
pve.py	PveClient (built on proxmoxer), build_client()	section 3.5
topology.py	build_topology(): the disk/storage/group join, D, S, U ^{sex} t, (C2) pins	sections 3.5-3.7, 5.1, 5.3 (C2)
reserve.py	compute_reserve_status(): (C4)/(C5), shared by show-load today and the solver later	section 5.3 (C4)/(C5)
loadmodel.py	compute_group_load(): the raw-series-to-l_d blend, min_coverage rejection, current L_s/u_s	section 4
gates.py	evaluate_group_gates(): reserve override, drift, imbalance — the act/no-act verdict, with reasoning	section 6
heuristic.py	run_heuristic(): seed/repair/descend, and evaluate_assignment(), the section 5.4 objective shared with the (unwritten) MILP path	sections 5.4/5.5
cli.py	Argument parsing, command dispatch, --manual, the mode-override rule, show-load, verify-storages	section 11.3

Not yet written: optimize.py, payback.py, schedule.py, execute.py — and, within modules that do exist, cooldown handling in gates.py (needs state.json’s timestamps; see [80-gates.md](#)), heuristic step 4 “polish” and (C2) format-compatibility eligibility in heuristic.py (see [90-heuristic.md](#)).

Why config.py depends on forecast.py

config.py’s semantic validation (`_check_forecast_window`) must reject a `window.lookback` too short for the configured `forecast.model` at startup (section 11.1) — that check needs `forecast.required_range_seconds()`. Since `forecast.py` in turn needs `config.ForecastConfig` for its type signatures, the import is deferred to inside the validation function rather than at module level, breaking what would otherwise be an import cycle. See the comment at the top of `config.py`’s `_check_forecast_window` if you touch this.

Where to read next

- [10-configuration.md](#) — how a YAML file becomes a validated Config, and where every default actually lives.
- [20-forecasting.md](#) — the forecaster protocol and its three implementations.
- [30-metrics.md](#) — the Prometheus client and the six `verify-metrics` checks.
- [40-cli-and-logging.md](#) — command dispatch, the `--mode` escalation rule, and why logs go to `stderr`.

- [50-pve-api.md](#) — the PVE API client, why it is built on proxmoxer, and two things verified against a real cluster.
- [60-topology.md](#) — the disk/storage/group join, the section 5.1.1 foreign-volume accounting, and the shared (C4)/(C5) evaluator.
- [70-loadmodel.md](#) — the section 4 raw-series-to-*l_d* blend, min_coverage rejection, and why tpmstate0/unusedN are exempt from it but efidisk0 is not.
- [80-gates.md](#) — the section 6 act/no-act verdict, why show-load’s gate line is diagnostic rather than a real decision today, and why cooldowns are not in this module.
- [90-heuristic.md](#) — the section 5.4/5.5 objective and the seed/repair/descend search, cross-checked against section 14’s exact objective totals, and what “polish” and format eligibility still owe.

How a YAML file becomes a validated Config

What does this page answer? What exactly happens between `pve-storage-drs -c foo.yaml plan` and a type-checked Config object, and where does each of the section 11.1 validation rules actually live? Describes `proxmox_storage_drs/config.py` and `config_schema.json`.

Resolution order

`resolve_config_path()` implements `IMPLEMENTATION_PLAN.md` section 11’s three-source order: `--config (cli_path)`, then `$PVE_STORAGE_DRS_CONFIG`, then `config.DEFAULT_CONFIG_PATH (/etc/pve/drs.yaml)`. It also returns `was_explicit`, which `load_config()` uses to decide the error message on a read failure: an explicitly named path that cannot be read is always fatal with no fallback, per section 11, but the *message* additionally points at `config/drs.example.yaml` only when the failing path was the unnamed default — there is nothing useful to suggest when the operator named the path themselves.

The two-stage validation

1. `_validate_schema()` loads `config_schema.json` via `importlib.resources` (it ships as package data — see `pyproject.toml`’s `[tool.setuptools.package-data]`) and runs it through `jsonschema`. Every object in the schema sets `additionalProperties: false`: a typo’d key is caught here as a validation error, not silently ignored, which is what keeps section 15.1’s knob-to-formula table trustworthy (a key with no formula would otherwise go unnoticed).
2. `_build_config()` walks the now-schema-valid raw dict into the frozen dataclasses defined earlier in the module, applying every default inline via `.get(key, default)`. These defaults are also what `config/drs.example.yaml` documents and what the manual’s [../manual/10-configuration.md](#) describes — there is deliberately no second copy of a default anywhere else in the codebase.
3. `_validate_semantics()` runs the section 11.1 rules `jsonschema` cannot express — cross-field comparisons, or rules needing a value derived from two independently-optional settings. Each rule is a small `_check_*` function appending to a shared `errors/warnings` list (kept separate functions rather than one large one, partly for flake8’s max-complexity and partly because each is independently testable). `errors` become one `ConfigError` naming every problem found, not just the first; `warnings` are returned to the caller (`ResolvedConfig.warnings`) for `cli.py` to log — section 11.1’s `saturation_load` check is the one warning-only rule today.

Every duration or size field is parsed **once**, at this point, via `units.parse_duration_seconds/parse_size_bytes`, and stored on the dataclass already in its canonical unit (seconds, bytes) — nothing downstream re-parses a string.

Secrets from the environment

`_build_config()` fills `proxmox.auth.password/token_secret` from `PVE_PASSWORD/PVE_TOKEN_SECRET` **only when the file’s own value is null** — a secret actually written in the file is used as-is, never silently overridden by an environment variable a deploy script forgot to unset. This one-directional fallback

is what section 11’s “keep secrets out of /etc/pve” guidance is for: the file can omit the field entirely and rely on the environment, or set it explicitly and take responsibility for the file’s own permissions.

What is deliberately *not* checked here

Two section 11.1 rules need live cluster data this module does not have and are therefore **not** part of `config.py`:

- storage ids actually existing in the cluster;
- `execution.source_release.timeout` / `gates.cooldown_per_storage` against a storage’s real `saferemove_throughput`.

Both belong to `pve.py/topology.py` (not yet written) and `pve-storage-drs verify-storages`, once they exist — see [../manual/30-safety-and-status.md](#) for current status.

ResolvedConfig: what the caller actually gets

`load_config()` returns a `ResolvedConfig`, not a bare `Config`: the resolved path and the file’s sha256 travel with it, because `IMPLEMENTATION_PLAN.md` section 11 requires every run to log both — a plan must always be traceable to the exact file that produced it, even though the file is never re-read after startup.

Forecasting: the protocol and its three implementations

What does this page answer? How does `forecast.py` decide how much history a model needs, and what does each of `quantile/seasonal_naive/holt_winters` actually compute? Describes `proxmox_storage_drs/forecast.py`.

The one rule, in one place

`IMPLEMENTATION_PLAN.md` section 10.1’s table says how much history each model needs, independent of `window.lookback` (the *decision* window). That rule has exactly one implementation: three small pure functions, `_quantile_required_range_seconds`, `_seasonal_naive_required_range_seconds` and `_holt_winters_required_range_seconds`, each called from **two** places — the free function `required_range_seconds()` that `config.py`’s startup validation uses (before any `Forecaster` is constructed), and the matching `Forecaster.required_range()` method on the concrete class. Duplicating the arithmetic between those two call sites, rather than sharing the pure function, is exactly the kind of “second copy of a rule” `AGENTS.md` section 5 forbids — a change to the Holt-Winters requirement made in only one of them would silently desynchronize config validation from what the forecaster itself believes it needs.

Why the upper bound, never the point estimate

Every `Forecaster.predict()` returns a `Forecast(point_estimate, upper_bound)`. Nothing in this module enforces that callers use `upper_bound` — that discipline belongs to the caller (the optimizer and the section 7.3 saturation guard, neither written yet). The asymmetry is stated in the module docstring and repeated here because it is easy to get backwards under time pressure: overestimating costs a slightly worse balance, underestimating risks scheduling a mirror onto a storage that is about to saturate.

QuantileForecaster

The default. `predict()` takes the `window.quantile/window.upper_quantile` percentiles of the raw series with no fitting at all. `_quantile()` implements linear-interpolation percentile matching `numpy.percentile`’s default method, by hand — deliberately not using `numpy`, since the tool must stay usable with only `requests + ruamel.yaml + jsonschema` installed (`IMPLEMENTATION_PLAN.md` section 2.1). `horizon` is accepted (to satisfy the `Forecaster` protocol) and immediately discarded: the quantile model has no notion of a horizon-dependent forecast.

SeasonalNaiveForecaster

Groups historical samples by hour-of-day (`int((ts // 3600) % 24)`) and takes the median as the point estimate, the configured upper quantile across the same bucket as the bound. `now_epoch_seconds` selects which hour-of-day bucket the *current* moment falls into; `predict()` ignores everything outside that bucket. A bucket with no samples returns `Forecast(0.0, 0.0)` rather than raising, matching `IMPLEMENTATION_PLAN.md` section 4's rule for an idle group's raw quantities: no data is zero, not an error, at this layer (the caller decides whether zero is trustworthy).

HoltWintersForecaster

`statsmodels` is an optional dependency (Suggests, not Depends — see `debian/control`), so it is imported **only inside `predict()`**, never at module level; `debian/tests/import-all` is what would catch a regression here. Two situations fall back to `QuantileForecaster`, both logged at warning level rather than failing silently:

1. fewer than `2 * seasonal_periods` samples in the series (section 10.2's own rule — a fit on too little data is worse than no fit);
2. `statsmodels` is not importable at all.

The upper bound is `point_estimate + residual_z * stdev(in-sample residuals)` — not derived from the model's own confidence interval, which `statsmodels's ExponentialSmoothing.fit()` does not expose directly for every configuration. `residual_z` (`forecast.holt_winters.residual_z`, default 2.0) is what an operator tunes if this bound turns out too tight or too loose in practice; there is no plan-mandated backtesting gate on it yet (`IMPLEMENTATION_PLAN.md` section 10.2's backtest-validation requirement is not implemented — tracked for the dedicated forecasting phase, section 12 phase 9).

build_forecaster()

The one factory function that turns a `ForecastConfig` plus the ambient window/metrics values into a live `Forecaster` instance. Nothing else in the codebase should construct a `QuantileForecaster/SeasonalNaiveForecaster/HoltWintersForecaster` directly once a caller exists that needs one from config — route it through here so the model-name-to-class mapping stays in one place.

The Prometheus client and verify-metrics

What does this page answer? How does `metrics.py` talk to Prometheus without ever touching the network in a test, and what does each of the six `verify-metrics` checks actually query? Describes `proxmox_storage_drs/metrics.py`.

`_SessionLike / _ResponseLike`: why not `requests.Session` directly

`PrometheusClient.__init__` accepts `session: _SessionLike | None`, a `Protocol` defined in this module with only the one `get(...)` method the client actually calls — not `requests.Session` itself. `.agents/testing.md` forbids any test talking to a real Prometheus; a structural protocol lets a test hand in a small dataclass with a matching `get()` instead of mocking or subclassing the real (network-capable) session. `PrometheusClient()` with no session argument still constructs a real `requests.Session()`, which satisfies the protocol structurally — nothing special is needed to make the real thing fit.

`_get()`: one request path, three endpoint shapes

Every one of `instant_query`, `range_query` and `label_values` goes through the private `_get()`, which does the HTTP call, checks the status code, parses JSON, checks Prometheus's own status: "success" field, and returns `payload["data"]`. That field's *shape* genuinely differs by endpoint — a dict with a result list for `query/query_range`, a bare list for `label/<name>/values` — which is why `_get()` is typed `Any` rather

than picking one shape; each public method knows which shape it asked for and narrows accordingly. Every failure mode (a transport error, a non-200 status, unparseable JSON, an unsuccessful Prometheus response) raises `MetricsError` from exactly this one place, so a caller catching `MetricsError` around any of the three public methods sees every failure mode uniformly.

PromQL construction

`build_rate_promql()` and `build_quantile_over_time_promql()` are pure string-building functions, deliberately factored out of any query-issuing code so they are unit-testable without a fake session at all — see `IMPLEMENTATION_PLAN.md` section 3.4 for the exact expressions they implement. `raw_metric_name()` is the one place that maps a `RAW_METRIC_FIELDS` entry ("read_ops", ...) to the configured metric name on a `MetricsConfig`, via `getattr` — this is what lets `verify_metrics()` iterate “every configured raw metric” without hand-listing the six field names a second time anywhere in this module.

`verify_metrics()`: six checks, six functions

Each `_check_*` function implements exactly one `IMPLEMENTATION_PLAN.md` section 3.3 numbered check and returns `Findings` (plus, where relevant, the data the report carries forward):

Function	Section 3.3 step	Returns
<code>_check_metric_names_exist</code>	1	findings only
<code>_check_sample_series</code>	2 + 3 (labels present)	findings, {metric_name: sample_labels}
<code>_check_device_label_collision</code>	4	one Finding or None (pure, no I/O)
<code>_check_coverage</code>	5	findings, {DiskKey: coverage_fraction}
<code>_check_observed_spacing</code>	6	findings, observed_spacing_seconds

`_check_coverage` and `_check_observed_spacing` both use `metrics.read_ops` as the one representative metric rather than probing all six: a coverage or spacing gap is a property of the underlying Telegraf scrape, not of which of the six raw quantities is read, so probing all six would be six times the Prometheus load for no additional information.

`VerifyMetricsReport.ok` is `True` iff no Finding has `level == "error"` — `cli.py`’s `verify-metrics` handler uses exactly this property to decide the process exit code, so there is one definition of “the verification passed.”

What `cli.py` does with the report

`_handle_verify_metrics` in `cli.py` constructs a `PrometheusClient` from `resolved.config.prometheus`, calls `verify_metrics()`, and renders either the human form ([level] message lines) or, under `--json`, a JSON-serializable dict — `DiskKey` keys are turned into "vmid:device" strings there, since JSON object keys must be strings and this module keeps no opinion about output formatting.

Command dispatch, the mode-override rule, and why logs go to `stderr`

What does this page answer? How does `cli.py` turn `argv` into a dispatched command, what exactly does overriding `--mode log`, and why does structured logging go to `stderr` instead of the `stdout` the plan originally specified? Describes `proxmox_storage_drs/cli.py` and `proxmox_storage_drs/logging_setup.py`.

One parser, global options first

`build_parser()` defines every global option (`-c/--config`, `--group`, `--mode`, `--json`, `-v/--quiet`, `--version`, `--manual`) on the **top-level** `ArgumentParser`, not on the subparsers, which is what makes

argparse itself enforce IMPLEMENTATION_PLAN.md section 11.3’s “global options before the subcommand” for free — there is no hand-written ordering check anywhere. `--help`’s defaults come from the real `default=` values passed to `add_argument`, combined with `argparse.ArgumentDefaultsHelpFormatter`; there is deliberately no second, hand-maintained list of options and defaults anywhere in this module, the manpage, or the manual (AGENTS.md section 8.5).

`main()` handles `--version` and `--manual/help` **before** requiring or loading any configuration — printing the version or the manual page must never depend on `/etc/pve/drs.yaml` existing. Only once a real subcommand is present does it call `config.load_config()`.

Command handlers: a dict, not a chain of ifs

`_COMMAND_HANDLERS`: `dict[str, CommandHandler]` maps each subcommand name to a function of `(ResolvedConfig, argparse.Namespace, effective_mode) -> int`. Every subcommand not yet implemented (`plan`, `apply`, `show-load`, `explain`, `verify-storages` — see [../manual/30-safety-and-status.md](#) for which ones that currently is) gets a handler from `_make_not_yet_implemented_handler(name)`, a closure factory rather than a lambda with a default-argument trick, because mypy’s strict mode cannot infer a bare lambda’s parameter types cleanly here — see the comment at that function if you are tempted to simplify it back to a one-liner. `_COMMAND_HANDLERS["verify-metrics"]` is overwritten with the real `_handle_verify_metrics` once `metrics.py` exists to back it. Adding a new implemented command is exactly this: write the handler, assign it into the dict, done — `main()`’s dispatch does not change.

The --mode escalation rule

`apply_mode_override(configured_mode, override)` is the entire implementation of section 11.3’s rule: moving *down* the ordering `dry-run < confirm < auto` (toward more caution) logs at `INFO`; moving *up* it (toward less caution — removing a barrier the operator’s own config set) logs at `WARNING`, naming both values. This is deliberately a small pure-ish function (its only side effect is the log call) rather than inlined into `main()`, so it is unit-testable against a `caplog` fixture without constructing a full CLI invocation.

Why logs go to stderr, not stdout

IMPLEMENTATION_PLAN.md section 2.1 originally specified `stdout` for structured JSON logs. This was amended in the same commit that added `logging_setup.py`: `--json` (section 9.5) emits the plan report on **stdout**, and interleaving log lines with that report on the same stream would make it unparseable by anything downstream. Logging instead goes to **stderr**; a `systemd` service unit captures both streams into the same journal regardless, so nothing about “captured by journald” is lost. This is the one place in the codebase where the implementation and IMPLEMENTATION_PLAN.md disagree with the *original* plan text on purpose — AGENTS.md section 7 rule 2 is why the plan text itself was corrected in the same commit rather than left to silently diverge from the code.

`configure_logging()` is called exactly once, from `main()`, after the `--version/--manual` early exits (which must not touch logging configuration at all) and before any config-dependent code runs, so that a config-loading failure is itself logged consistently with everything after it.

JsonFormatter: what ends up in one log line

Every field on a `logging.LogRecord` that is *not* one of the standard attributes (computed once, at import time, into `_STANDARD_RECORD_ATTRS`) is folded into the JSON payload — this is what lets any module call `logger.info(..., extra={"event": "gate_decision", "threshold": 0.2})` and have `event/threshold` appear as top-level JSON keys with no formatter changes required. `sort_keys=True` keeps output byte-stable for anything that diffs log lines in a test or a log-aggregation pipeline.

--manual: preferring man(1), falling back only when necessary

`show_manual()` tries `man pve-storage-drs` first, via `subprocess.run` rather than `os.execvp`: `man(1)` itself detects a non-tty stdout and disables its pager automatically, which is what satisfies “never answer with a URL alone” (AGENTS.md section 8.5) whether the invocation is interactive or piped — there is no need for this module to duplicate that tty detection. `_fallback_manual_text()` is reached only when `man(1)` itself is missing or has no entry (an uninstalled source checkout, or a minimal system with no `man-db`): it walks up from `cli.py`’s own file location looking for `man/pve-storage-drs.1.md` in a source tree, and only if that also fails prints a short message naming the real installed-package path and the in-tree file — a local, actionable path, never a bare URL.

The Proxmox VE API client

What does this page answer? Why is `pve.py` built on `proxmoxer` instead of a hand-rolled ticket/CSRF client, and what does each of its methods actually call? Describes `proxmox_storage_drs/pve.py`.

Why proxmoxer

IMPLEMENTATION_PLAN.md section 3.5 explains the decision in full; the short version is `proxmoxer`’s **backend abstraction**. The same `ProxmoxAPI` attribute-chaining interface (`api.nodes(node).qemu(vmid).config.get()`) is available over plain HTTPS (what `build_client()` uses today) or over SSH, either shelling out to the system’s own `ssh+pvesh` (`openssh`) or with an in-process client (`ssh_paramiko`). A deployment that cannot open the API port to the management host becomes a different keyword argument in `build_client()`, with no change anywhere else — not to `PveClient`’s methods, and not to any of their callers.

build_client(): the one place auth is assembled

`config.py`’s `ProxmoxConfig` keeps `verify_ssl` (bool) and `ca_file` (path-or-None) as two separate knobs, because that is how an operator thinks about them; requests (and so `proxmoxer`’s https backend) wants one value, a bool or a CA-bundle path, passed as `verify`. `build_client()` is the one place that reconciles the two. It is also the one place `proxmox.auth.token_id`’s `user@realm!tokenname` format is split into `proxmoxer`’s separate `user/token_name` keyword arguments — never duplicate that parsing anywhere else.

Every failure `ProxmoxAPI(...)` itself can raise (`proxmoxer.AuthenticationError`, or a `requests.RequestException` if the host is simply unreachable) is caught here and re-raised as `PveApiError`, so every caller of `build_client()` and every `PveClient` method has exactly one exception type to catch.

PveClient: one method per section 3.5 endpoint

Every method is a single API call, wrapped through the private `_call()` helper, which is the one place `proxmoxer.ResourceException` (HTTP-level API errors) and `requests.RequestException` (transport failures `proxmoxer`’s https backend does not wrap itself) both become `PveApiError`. No method caches anything, and the per-run topology cache belongs to `topology.py` (section 3.5), which is also where the bounded thread pool for the O(VMs) config fetch lives (60-topology.md, REVIEW.md P-02).

`_call()` does retry exactly one failure mode: `proxmoxer.AuthenticationError` gets one reauthenticate-and-retry cycle (REVIEW.md P-01) before becoming a `PveApiError`, via a reauthenticate callback `build_client()` wires in — a full fresh login through `_build_api()` again, not a reuse of the rejected ticket. This is not a general retry policy (a `ResourceException` or a transport failure still fails immediately, on the first attempt); it exists specifically because a ticket can expire for reasons with nothing to do with the call that hits it — a long `apply --confirm` wait, a suspended process, a clock jump — and `proxmoxer`’s own lazy per-request renewal only notices the age of its own clock, not whether the server-side ticket actually outlived

a gap that long. A test double built with a bare `PveClient(fake_api)` (no `reauthenticate=`) gets none of this — it behaves exactly as it did before P-01, which is what every existing fake in `tests/unit/fakes.py` relies on.

`build_client()` also applies `proxmox.ticket_refresh_seconds` to the constructed session, best-effort: proxmoxer 2.x has no constructor argument for a password/ticket auth's refresh interval (it is a hard-coded `renew_age = 3600` class attribute on `ProxmoxHTTPAuth`, checked lazily on whichever request happens to run after it elapses), so `_apply_ticket_refresh_seconds()` reaches into `api._backend.auth` — undocumented proxmoxer internals, not its public interface — to override that instance's `renew_age` after login. Deliberately narrow: it only ever *tightens* proxmoxer's own schedule, never replaces its refresh logic, and if a future proxmoxer version reshapes this (or the backend is not `https`, or the auth is API-token, which has no ticket at all), the `getattr/hasattr` guards make it a silent no-op — the P-01 `reauthenticate-and-retry` above still catches the case this can't reach.

`storage_definitions()` deliberately uses `GET /storage` (the list form), never the singular `GET /storage/{id}`, even though only one storage's config is needed in some callers. This was verified empirically against a real PVE 9.2.11 cluster during development ([[dev-cluster-access]] in the maintainer's notes, not reproduced here): an API token granted only `Datastore.Audit` can list every storage's full config — including `saferemove/saferemove_throughput` — via `GET /storage`, but the identical data via `GET /storage/{id}` for the same storage returns 403 Forbidden (`Datastore.Allocate`). The single-item route apparently backs an edit-UI flow gated by the ability to change the config, not merely read it. Since the list form returns the same per-storage object for every storage in one call, there is no reason to ever call the singular form, and using it would quietly force a Storage DRS token to hold more privilege than the tool actually needs. IMPLEMENTATION_PLAN.md section 3.5 was corrected to match once this was found — see that section's note.

`storage_content()` returned an empty list — not an error — on every storage, node and content-type filter tried, until the token also held `Datastore.Allocate` on that storage. This was genuinely surprising and worth calling out precisely because it is a *silent* false negative: `Datastore.Audit` alone gives a clean 200 with `{"data": []}`, which looks exactly like “this storage legitimately has no tracked content” rather than “the caller cannot see it.” It was not guessed at — see IMPLEMENTATION_PLAN.md section 3.5's note, which records the exact before/after (0 entries with `Audit` only, 37 and 1 entries respectively for two real storages once `Datastore.Allocate` was added) on the same live cluster. **The practical consequence: `pve.py`'s own “keep the token to Audit-only” goal for `storage_definitions()` does not extend to `storage_content()` — a Storage DRS token genuinely needs `Datastore.Allocate` (bundled with `Datastore.Audit` in one role) on every storage it manages, or `topology.py` will silently compute U^{next} (section 5.1.1) as if every storage were empty of untracked volumes, with no error to notice by. docs/manual/00-installation.md states the exact minimum ACL grant an operator needs; this is not a place to economize on privilege out of a preference that turned out not to match reality.**

Two more things about Proxmox's ACL model this surfaced, both now reflected in the manual's token-setup instructions: **an API token's effective permission is the intersection of its owning user's permissions and whatever is granted to the token itself** when privilege separation is on (the default) — granting a role to only one of the two has no effect, both need it; and the grant that was confirmed to work was made explicitly at each `/storage/{id}` rather than only at the parent `/storage` path, so that is what the manual instructs, rather than relying on propagation that was not cleanly isolated as working or not.

move_disk(): the one and only bytes/s -> KiB/s conversion

Section 9.2 is explicit that `bwlimit`'s bytes/s-to-KiB/s conversion happens “at the call site and nowhere else.” `PveClient.move_disk()` is that call site: every caller in this codebase, present and future, passes `bwlimit_bytes_per_sec` in the same unit `migration.bwlimit_bytes_per_sec` already uses, and the `divi-`

sion by 1024 (rounded, since the API wants an integer) happens exactly once, inside this method. `format` defaults to `None` and is only ever forwarded when a caller explicitly passes one — `move_disk` without `format=` preserves the source format, which is what section 3.5 requires by default.

Building the disk/storage/group model

What does this page answer? How does `topology.py` turn `pve.py`'s raw API responses into the per-group D/S sets the rest of the engine needs, and where does each `IMPLEMENTATION_PLAN.md` section 5.3 (C2) pin condition actually get decided? Describes `proxmox_storage_drs/topology.py` and `proxmox_storage_drs/reserve.py`.

D means every disk this module places, pinned or not

This is stated at the top of `topology.py`'s own docstring because it is the single easiest thing to get backwards: D (section 5.1) is *every* disk this module puts into a group's `Group.disks`, whether or not (C2) pins its placement. A disk this module never sees at all — a stopped VM excluded by `exclude.running_only`, or a disk whose current storage is not in any configured group — is what “foreign” (`Useext`, section 5.1.1) means. Config-excluded disks (`exclude.vmid`s/`exclude.disks`/`tags`) are *not* foreign: (C2) pins them into D specifically so their bytes still count in (C4)/(C5) and their fragmentation toward κ (section 3.6, “Pinned disks are modelled, not ignored”). An earlier draft of section 5.1.1 listed config-excluded disks as foreign, which directly contradicted (C2) — that self-contradiction was found and fixed in the same commit that first implemented this join; see that section's note if you need the history.

One pass, in the order section 3.5 lists

`build_topology()` fetches everything it needs exactly once, in five stages (`_fetch_cluster_data`, then a loop over `client.vm_resources()` calling `_collect_vm_disks` per VM, then `_build_storages`):

1. `storage_definitions()` (the list form — see 50-pve-api.md), validated against the configured groups (`_validate_group_storages`): a storage that does not exist, or that lacks images in its content types, is a fatal `TopologyError` (section 11.1: “storage ids exist in the cluster”). A storage that exists but is not marked shared is a warning, not a fatal error — plausible, if unusual, for a single-node group.
2. `storage_resources()` once, to pick one active node per storage (`_pick_active_node`, preferring one reporting status: “available”).
3. `storage_content()` and `storage_status()` per group storage.
4. `vm_config()` and `vm_snapshots()` per considered VM. A VM is *not* fetched at all when `exclude.running_only` is set and it is stopped — the one case `cluster/resources`'s own fields can decide without a config fetch. Every other VM is fetched, including config-excluded ones: whether a specific disk turns out to be “ungrouped” is only knowable after seeing which storage it is actually on, and a config-excluded VM's disk still needs its size for (C4)/(C5) (previous section).

This is the one step that runs across a bounded thread pool (`config.proxmox.read_workers`, REVIEW.md P-02): section 3.5's whole point is that `GET .../qemu/{vmid}/config` has no batch form, so for a several-hundred-VM cluster this is the read phase concurrency actually helps. It is split into `_fetch_vm()` (network I/O only, run by the pool, one call per considered VM) and `_join_vm_disks()` (pure — the same function the old sequential `_collect_vm_disks()` was renamed from, unchanged in what it computes). `ThreadPoolExecutor.map()` returns each VM's fetch in the same order `client.vm_resources()` listed it, even though the fetches themselves complete in whatever order the pool schedules them, so the join phase runs single-threaded and in the original order — `disks_by_group`, `referenced_volid`s and warnings end up byte-for-byte identical to a fully sequential run regardless of `read_workers` or of thread scheduling, and never need a lock. `PveClient`

itself may reauthenticate mid-fetch if several threads hit an expired ticket at once (50-pve-api.md's P-01 note); that reauthentication is what its own lock serializes, not this join.

5. Non-QEMU resources (type != "qemu") are skipped outright — this tool never touches LXC containers, and section 3.5's read/write paths are qemu-only throughout.

The lock field comes from vm_config(), not a separate call

Section 9.3's own pseudocode says lock is readable from either /qemu/{vmid}/config or /status/current. _collect_vm_disks reads it from the config response already being fetched for the disk join, so planning-time lock detection costs no extra API call. /status/current is still needed later, but only immediately before a specific move (section 9.2's pre-move re-validation, execute.py, not yet written) — never here.

Pin priority: _pin_reason()

One function, checked in the fixed order section 5.3 (C2) lists its conditions: config exclusion (VM-level, then exclude.disks), a real snapshot or an unreferenced companion volume (section 3.7, _disk_snapshot_or_orphan_reason — the "current" entry GET .../snapshot always returns, verified on a live cluster, is filtered out before counting), a VM lock, then an unusedN disk when exclude.include_unused_disks is false. A disk gets at most one reason; the first that applies wins, matching how an operator would explain it.

Sizes: content is authoritative, config is the fallback

_resolve_disk_size_and_format looks up the volume in that storage's content listing first (section 3.5: authoritative for what is actually allocated) and falls back to parsing the VM config's own size= only when the volume is missing from content — logged as a warning naming the disk, since assume-thick-provisioning sizing from a stale or unlisted config value is a real, if rare, source of drift. Two independently verified reasons this fallback path is not merely defensive: the Datastore.Allocate-vs-Audit privilege gap (50-pve-api.md), and an unusedN volume, deleted directly on the storage backend outside Proxmox (confirmed with the cluster's operator), that PVE's own config still referenced and that PVE itself never noticed or refused to boot the VM over — IMPLEMENTATION_PLAN.md section 3.6's note. _parse_pve_config_size_bytes parses PVE's own config-file size suffixes (512G meaning binary GiB, no explicit i) — deliberately not units.py's parser, which is for *this project's* config file, a different and coincidentally similarly-shaped format.

U_{set}: everything not referenced

Every disk actually placed into a group's D records its valid in referenced_valids[storage_id] as it is processed. _build_storages then sums the size of every content entry on that storage whose valid was never referenced — orphans, templates, other groups' foreign volumes, and disks of VMs this run never fetched (stopped-and-excluded, or ungrouped) all fall out of this one computation with no special-casing per category, which is exactly section 5.1.1's definition read literally.

reserve.py: one (C4)/(C5) evaluator, shared

compute_reserve_status() is deliberately its own module, not a method on Storage: the solver (optimize.py/heuristic.py, not yet written) will need to evaluate the identical formula against a *candidate* assignment, not only the current one, and pve-storage-drs show-load's reporting needs it against the current assignment today. Both call the same function (AGENTS.md section 5) — show-load's use of it is already exercised end-to-end (see cli.py's _render_show_load_human/_json). The section 14 worked example's initial state (tests/unit/test_reserve.py) is the proof this reproduces the plan's own arithmetic exactly, not merely a self-consistent unit test.

The load model

What does this page answer? How does `loadmodel.py` turn six raw Prometheus series into the single per-disk `l_d` section 4 defines, and what happens when the data for one disk cannot be trusted? Describes `proxmox_storage_drs/loadmodel.py`.

One function, one group, already-built input

`compute_group_load()` takes an already-built `topology.Group` (one group's D/S, section 5.1) and a `metrics.PrometheusClient` and returns a `GroupLoad`: one `DiskLoad` per disk (`l_d`) and one `StorageLoad` per storage (`l_s`, `u_s`), plus the group's `u*` and whether the whole group is idle. It never touches the PVE API and never builds topology itself — matching `reserve.py`'s identical “given the join, evaluate one formula” framing, not a coincidence: both are meant to be called again, unchanged, by the solver once it exists.

Six queries, not six-times-groups queries

The six raw quantities (section 3.4) are fetched with **one instant query each**, unfiltered by group — `sum by (vmid, device) (...)` already returns every disk Prometheus currently reports, and picking out one group's keys in Python is free, whereas six *more* Prometheus queries per group is not. `read_time_ns/write_time_ns` are converted from nanoseconds to seconds here, engine-side (`_combined_raw`), for the same reason F-03 moved `read_factor/write_factor` engine-side: one tested constant beats the same /1e9 repeated across deployed PromQL strings.

Known limitation, not a bug: a config with more than one group re-runs these same six unfiltered queries once per group — correct (each group's coverage/rejection decisions are independent) but wasteful for a many-group cluster. Fetching once and slicing per group would be a straightforward follow-up if a real deployment's group count ever makes this Prometheus load worth avoiding; nothing about `compute_group_load()`'s per-group signature forces the current one-fetch-per-call shape, it is just what the first implementation does.

Coverage rejection excludes a disk from the group total, not just from its own `l_d`

Section 3.4: “reject a disk whose sample coverage over `W` is below `window.min_coverage` and fall back to its last known load from `state.json`... never treat missing data as zero load.” Two things worth being explicit about, since the plan states the rule but not this mechanism:

1. **The rejected disk's raw contribution is excluded from `T_g/0_g/B_g` entirely**, not merely capped or zeroed in place — one noisy or half-missing series must not bias every *other* disk's normalized share. `l_d` for every accepted disk is computed from the accepted-only totals.
2. **The rejected disk's own `l_d` is substituted afterward**, from `last_known_loads` if the caller has one (keyed by `topology.Disk.key`), or flagged `0.0` if not. `DiskLoad.flagged_reason` is set either way, so a caller can never mistake a flagged `0.0` for a genuinely idle disk — `show-load` renders it as an explicit `Δ` line, never silently.

`state.json` persistence (the actual `last_known_loads` source) is not yet written (section 11.2) — `compute_group_load()` accepts the mapping today so this module needs no change when it exists; until then every call site passes nothing, and every coverage-rejected disk with no seed is flagged `0.0`, honestly.

tpmstate0/unusedN are never rejected; efidisk0 is

Section 3.4's own note: `tpmstate0` and `unusedN` are not QEMU block devices and legitimately emit no `blockstat` series — `l_d = 0` for them is correct, not a data-quality problem. `_is_metrics_expected_absent()` encodes exactly this device-name distinction so these two kinds are never flagged for low coverage, while `efidisk0` — a real QEMU drive that *should* appear — is held to the same `min_coverage` bar as `scsi/virtio/ide/sata`. Getting this backwards either way is a real bug: exempting `efidisk0` would

hide a genuine metrics gap on it; not exempting `tpmstate0/unusedN` would spam a flag on every run for every cluster, for a “gap” that isn’t one.

T_g = 0 is “idle”, computed only from accepted disks

`GroupLoad.idle` is `True` exactly when every *accepted* disk’s `raw_t` is zero — the guard section 4 requires (“if `T_g = 0` the entire group is idle, skip it”). An all-rejected group (e.g. a total Prometheus outage mid-window) also comes out `idle=True` by this same computation, which is not quite literally accurate (the truth is “unknown”, not “idle”) but is harmless: both mean “do not act”, which is the only decision that follows from either. A caller that wants to distinguish the two can — every disk’s `flagged_reason` is still there to check.

L_s/u_s are the *current* assignment, not a candidate one

`StorageLoad` sums `l_d` over disks whose `Disk.current_storage` already points at that storage — today’s state, exactly like `reserve.py`’s `compute_reserve_status()`. The solver (`optimize.py/heuristic.py`, not yet written) will evaluate the identical per-disk `l_d` values against a *candidate* assignment instead; nothing in `loadmodel.py` needs to change for that, since `l_d` itself does not depend on which storage a disk is currently on.

What is still missing from phase 3

`IMPLEMENTATION_PLAN.md` section 12 phase 3’s own “done when” is “correct act/no-act decision per group, with the reasoning shown” — that is the section 6 drift/imbalance gates, cooldowns, and the reserve override, none of which exist yet. `loadmodel.py` is the load computation those gates consume, not the gates themselves; `pve-storage-drs show-load` reports `l_d/L_s/u_s` today, but no command yet decides whether a group should be re-balanced.

Gating: deciding whether to act at all

What does this page answer? How does `gates.py` turn a `loadmodel.GroupLoad` and per-storage `reserve.ReserveStatus` into the section 6 act/no-act verdict, and why isn’t cooldown handling here? Describes `proxmox_storage_drs/gates.py`.

One function, three gates, in the order section 6 lists them

`evaluate_group_gates()` is pure — no network I/O, no `state.json` access — taking a group’s already-computed `GroupLoad`, a `storage id -> ReserveStatus` mapping (the caller’s job, via `reserve.compute_reserve_status()` — this module never recomputes it, per `AGENTS.md` section 5’s “one implementation of every rule”), the group’s `GatesConfig`, and the load vector recorded at the group’s last *executed* balance (`last_load`, `None` if there isn’t one). It checks, in section 6’s own order, stopping at the first that decides:

1. **Reserve override.** Any storage in `reserve_statuses` with `.violated` set forces `act=True` immediately, bypassing drift and imbalance entirely — section 13’s “safety is not subject to hysteresis” made literal, not just documented intent.
2. **Drift gate.** $\|l_{\text{now}} - l_{\text{last}}\|_1 / \|l_{\text{last}}\|_1 \geq \text{gates.drift_threshold}$, vectors aligned over the **union** of disk keys (a disk absent from one side contributes its full load in the other as drift — a new or deleted disk is a genuine change to the group’s I/O profile). `last_load=None` skips this gate outright, per section 6’s own degenerate-case table — not “treat as zero drift”, which would make an operator’s very first run fail to act on an already-imbalanced cluster.
3. **Imbalance gate.** $(\max_s u_s - \min_s u_s) / u^* \geq \text{gates.imbalance_threshold}$, u^* from `GroupLoad.average_utilization`. $u^* == 0$ (idle group, section 4) always means no-act — there is nothing to spread evenly.

last_load=None everywhere in this codebase, for now

state.json (section 11.2) is not yet written, so nothing in this codebase has an actual “last executed balance” to read. Every current caller (show-load) passes last_load=None, which per the degenerate-case table means the drift gate is always skipped and every decision reduces to “reserve override, else imbalance” — an accurate live snapshot, but not the real hysteresis behaviour section 6 describes once a plan actually executes and starts recording `l_last`. `evaluate_group_gates()` itself needs no change when state.json exists; only its caller does.

show-load’s gate line is diagnostic, not a real decision

pve-storage-drs show-load prints each group’s verdict (Group <name> → ACT: <reason> / → NO ACTION: <reason>) reusing `GateDecision.reason` verbatim — deliberately the single source of truth for the wording, rather than a second, similarly-but-not-identically phrased string living in `cli.py`. This is useful today (an operator can see whether the tool *would* act, and why) but is not what plan/apply (not yet written) will actually gate an execution on, once last_load is real: show-load’s verdict, precisely because it always passes last_load=None, can only ever show “reserve override” or “imbalance”, never a drift-suppressed “no action” a real run with history would correctly show.

Cooldowns are not this module’s job

`gates.cooldown_per_disk/cooldown_per_storage` decide which *individual* disks or storages are movable once a plan is already being built — the same kind of decision (C2)’s other pin reasons make in `topology.py`, not a group-wide act/no-act verdict. They need state.json’s per-disk/ per-storage “last moved at” timestamps, which do not exist yet either, so implementing them here now would be speculative machinery with no caller and no real timestamp source to test against. They belong with whichever module builds the solver’s movable-disk set (`heuristic.py/schedule.py`, phase 4) when that is written.

The heuristic solver

What does this page answer? How does `heuristic.py` turn a group’s loads into a target assignment, why is the section 5.4 objective its own function rather than baked into the search, and what does this module deliberately not do yet? Describes `proxmox_storage_drs/heuristic.py`.

evaluate_assignment() is the objective, kept separate from the search

Section 5.5 requires it explicitly: “the heuristic must use the **same** feasibility and objective functions as the MILP path so the two backends are directly comparable.” `evaluate_assignment()` takes a Group and an arbitrary candidate Assignment (dict[Disk.key, storage id], not necessarily reachable from the current state by any sequence of moves) and returns an `ObjectiveBreakdown` — the section 5.4 objective’s five terms kept separate, not collapsed into a single number, because `explain` (not yet written) needs the arithmetic shown, and because the tests cross-checking this against section 14’s worked example need to assert each term individually, not just a total that could be right for the wrong reasons.

This is the same design choice `reserve.py` made first: a pure function of *data*, with no notion of “the current run” baked in. `heuristic.py` reuses `reserve.compute_reserve_status()` directly for (C4)/(C5), passing it a `storage_of` callback closed over the candidate assignment being evaluated — the extension point `reserve.py`’s own docstring already anticipated (“the solver will pass a candidate assignment through the identical shape”) before this module existed to use it.

Repair is unconditional, not weight-driven

Section 13: “the reserve is never traded against balance.” The heuristic makes this literal rather than merely well-weighted: `_repair()` runs *before* `_descend()` and fixes any (C5) violation by moving disks off the worst-violating storage, one iteration at a time, regardless of what the objective weights say — there is no beta/gamma value large enough to switch this off. `objective.reserve_violation_penalty` still

appears in `ObjectiveBreakdown.reserve_penalty_term`, but only for *reporting* a residual that repair could not fix (physically impossible, not merely unattractive) — by the time `_descend()` runs, every reserve violation `_repair()` could resolve already has been, so in the common case this term is exactly zero and does not influence descend’s choices at all. `test_reserve_violation_is_repaired_even_with_beta_high_enough_to_forbid_balance_moves` is the regression test for this: `beta=100` (so no purely-cosmetic balance move could ever be worth a migration) still sees the repair move happen.

Descend explores swaps, not only single moves

Section 5.5 is explicit that swaps are not an optional refinement: “when both storages are near their capacity limit, no single move is feasible, and only an exchange of two disks can improve the balance.” `_descend()` evaluates every single-disk move *and* every pairwise swap of two movable disks each iteration, applying whichever most reduces the objective, and stopping when nothing does or after `heuristic_` iterations. Both are plain re-evaluations of `evaluate_assignment()` against a trial dictionary — no separate swap-feasibility logic exists, because a swapped assignment is just another candidate the same objective function scores. `test_descend_uses_a_swap_when_no_single_move_is_feasible` constructs exactly the capacity-locked scenario the plan describes (two storages each with headroom smaller than any single disk) and checks the only reachable improvement — a swap — is the one found.

objective.spread_metric: two different quantities, not one rescaled

Section 5.4/(C6) offers two imbalance forms, and `evaluate_assignment()` implements both: “`l1`” (default) is $\alpha * \sum(e_s)$, every storage’s absolute deviation from u^* summed; “`minmax`” is $\alpha * \max(u_s)$ — the plan’s own $t \geq u_s$ for all s , the raw utilization of the single hottest storage. These are not the same quantity with a different aggregation: $\max(u_s)$ and $\max(e_s)$ can disagree, because a storage sitting far below u^* produces a large deviation e_s that minmax was never meant to notice (the manual’s own words: “indifferent to a second nearly-as-bad storage” — nearly-as-bad on the *high* side; a cold storage is not what minmax was designed to react to at all). `ObjectiveBreakdown` carries both `spread_e` (deviations) and utilization (raw u_s) regardless of which metric is active, computed once in the same loop, so switching metrics never means computing something the other mode didn’t already have on hand.

REVIEW.md’s Q-01 found this config knob silently ignored by an earlier version of this module (always `L1`, `objective.spread_metric` never read) — the same class of gap P-01/P-02 found in `pve.py/topology.py`. `test_spread_metric_minmax_is_indifferent_to_a_second_nearly_as_bad_storage` is the regression test, built directly from the manual’s own claim: two scenarios with an identical hottest storage but a materially different second-worst one score identically under minmax and differently under `l1`.

Proof this reproduces the plan, not just itself

`tests/unit/test_heuristic.py` doesn’t stop at “the assignment matches the prose.” At the default weights it asserts `ObjectiveBreakdown.total` equals 2.533333 (three-move plan) and, at `beta_move_count: 0.50`, 3.158333 (two-move plan) — the exact totals REVIEW.md Appendix A independently re-derived by hand from the plan’s own numbers, not values this module invented and then asserted against itself. Getting both to five decimal places is strong evidence the objective’s five terms, their units (TiB for size, average in-flight I/O for load), and the search that picks among them are all correct together, not merely internally consistent.

What this pass deliberately does not do

- “Polish” (heuristic step 4) — reuniting a fragmented VM when doing so does not worsen imbalance beyond `gates.imbalance_threshold`. Not implemented: the section 14 fixture’s exact three-move and two-move solutions are both reachable by repair+descend alone (proven by the tests above), because the objective’s own κ term already makes descend prefer co-location whenever it is not too costly. Polish would matter for an affinity fix that needs three or more disks to rotate

simultaneously — outside single-move/pairwise-swap reach — which the one fixture that exists to validate this module does not exercise. A real gap, tracked here rather than silently absent.

- **(C2) format-compatibility eligibility** — `topology.Storage` does not yet carry the storage type/format information that rule needs (it is resolved internally in `topology.py`'s `_default_format` but never exposed on `Storage`), so every group storage is currently an eligible target for every movable disk. The section 14 fixture is homogeneous (three identically-typed storages) and does not exercise this gap either.
- **CP-SAT/CBC coefficient scaling (section 5.5)** — belongs to `optimize.py`, not yet written. `evaluate_assignment()`'s plain floating-point objective is what the MILP path will need to agree with, not a stand-in for its own separately-scaled integer objective.
- **Wiring into any CLI command.** `heuristic.py` is not yet called from `cli.py` — plan/explain/apply remain stubs (`docs/manual/30-safety-and-status.md`). The next piece is `schedule.py` (section 8: ordering a target assignment's moves under the transient reserve invariant), after which both can be wired into `plan`'s dry-run output together.